



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

LAB04: PULP-NN - Deep Neural Network Inference on PULP

Davide Nadalini – d.nadalini@unibo.it

Alberto Dequino – alberto.dequino@unibo.it

Francesco Conti – f.conti@unibo.it

Objective of the Class

Intro: PULP platform and the PULP-SDK

Tasks: some basics of C programming on PULP:

- Parallelization on the PULP architecture
- Matrix-multiplication
- Fully Connected layer with vectorized instructions
- 2Dconv
- profiling code execution

Programming Language: C

Lab duration: 3h

The class is meant to be interactive: coding together and on your own!

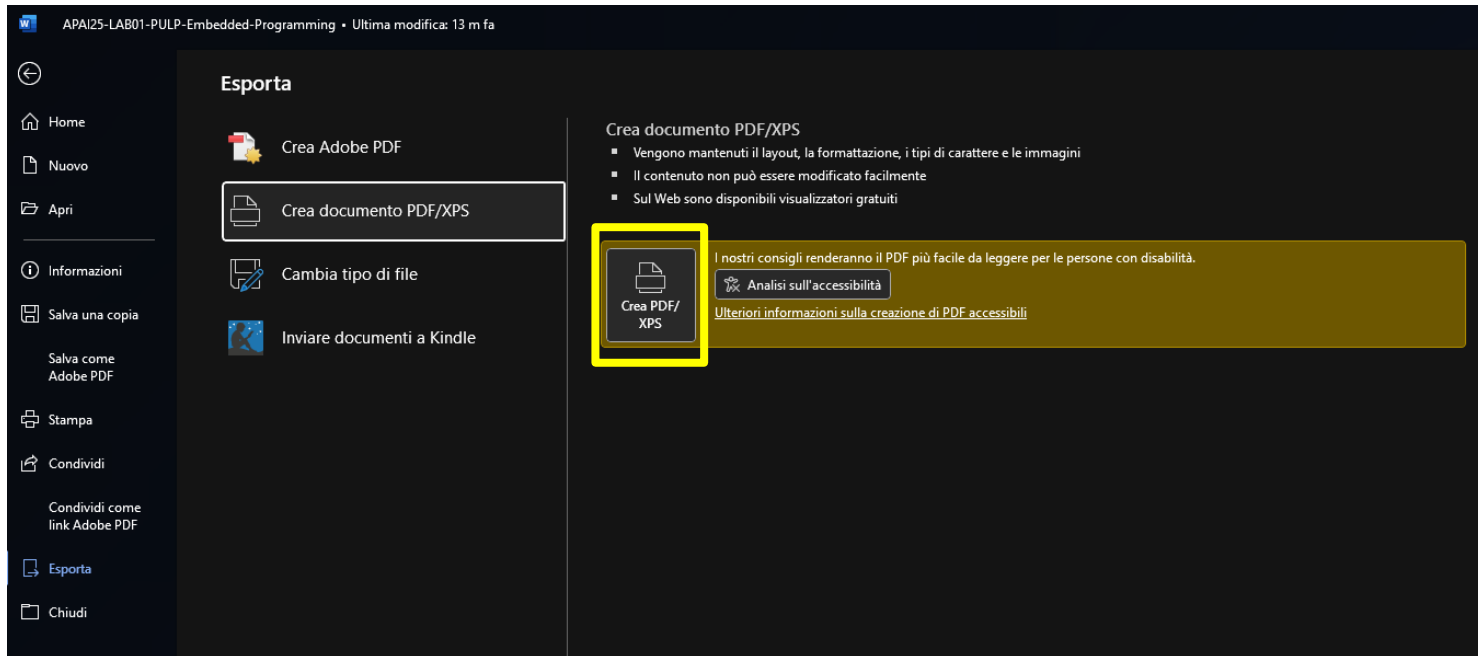
Deadline:

Oct 31st 2025

How to deliver the Assignment

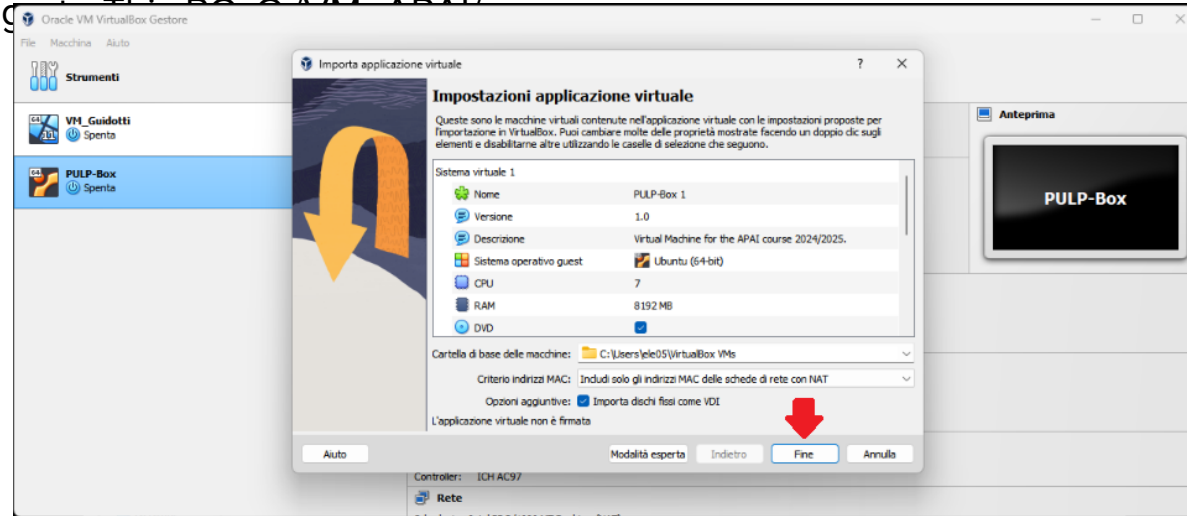

You will deliver ONLY the PDF assignment, no code

- Download the assignment file from Virtuale.
- Fill the results required by the assignment.
- Export to pdf format.
- Rename the file to: LAB<number_of_the_lesson>_APAI_<your_name>.pdf
- Use Virtuale platform to load ONLY your .pdf file



Opening the VM and VSCode

1. On the lab's PCs, open the file explorer and go to `C:\Users\lele\Downloads\`
2. Double click on PULP-box.ova
3. VirtualBox opens, just click on “Fine”
4. Wait for the VM to be imported
5. Open the VM with “Avvia”



Password is 'pulp'

Opening the Docker with VSCode

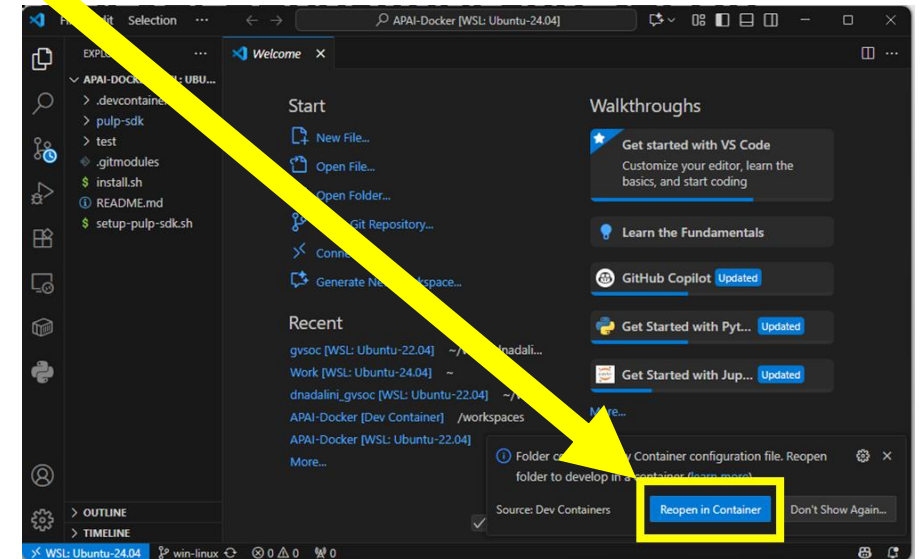
1. Open a terminal (right click – open a new terminal)
2. From the terminal, open VSCode in the folder of the Docker

```
$ cd APAI-Docker  
$ code .
```

1. Reopen the APAI-Docker folder in VSCode (click on “Reopen in container”)
1. Now you can use the **integrated terminal (open with CTRL+J)** to run your applications!

IMPORTANT: every time you open a **new terminal** to work on PULP, launch

```
$ source setup-pulp-sdk.sh
```



Getting Started:

IMPORTANT: activate the pulp-sdk module file every time a new shell is open.

```
$ source setup-pulp-sdk.sh
```

HOW TO RUN THE CODE:

```
$ git clone https://github.com/EEESlab/APAI25-LAB04-PULP-NN
$ cd APAI25-LAB04-PULP-NN
$ cd <folder_you_want>
$ make clean all run
```

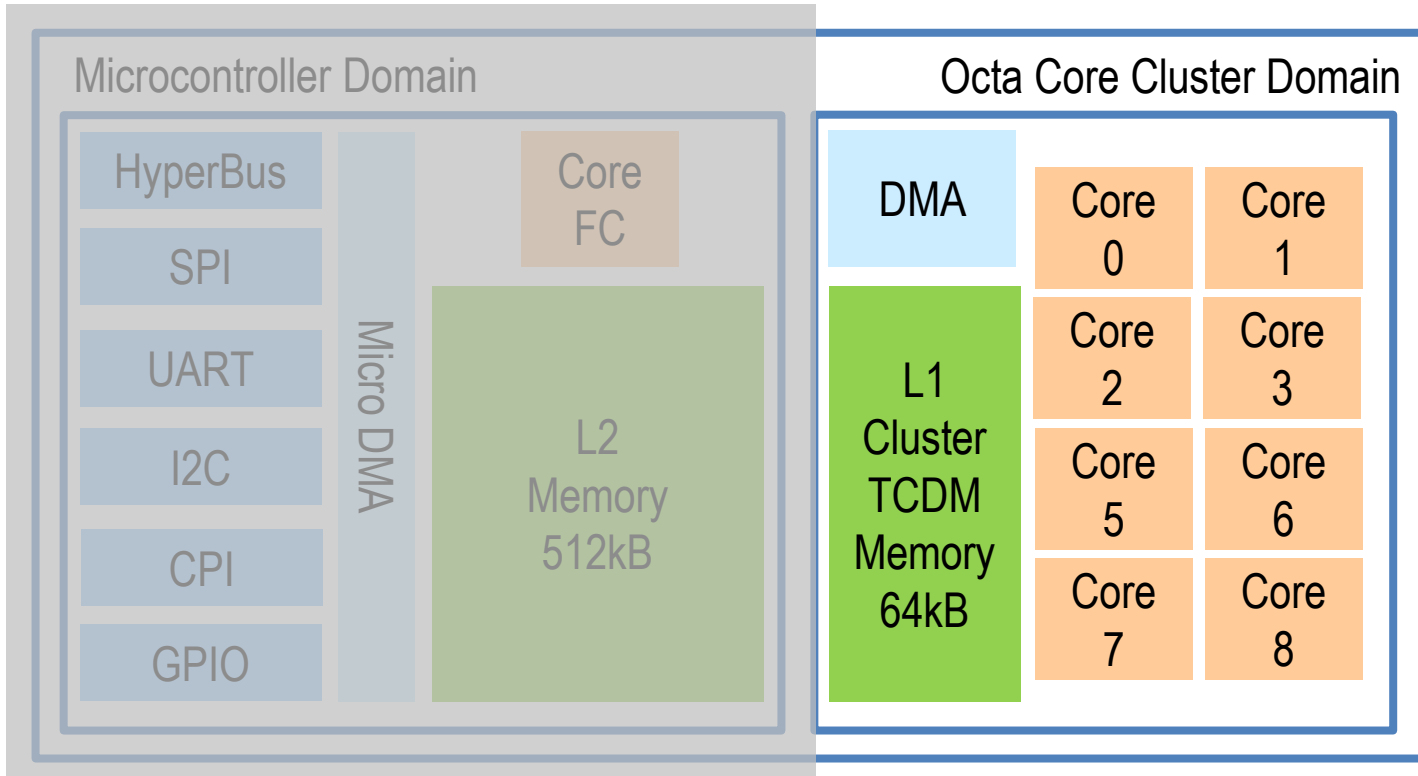




ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

INTRO

PULP Platform: today we focus on the 8-cores cluster



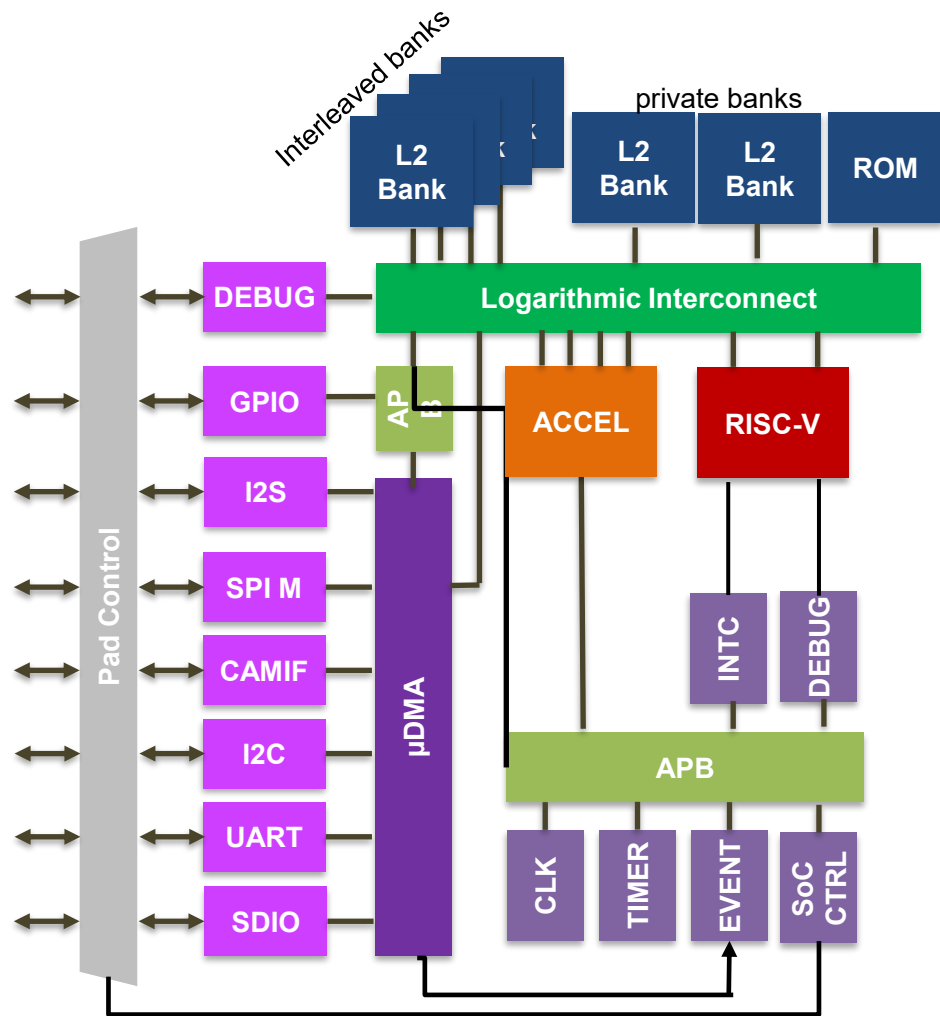
- **Cores:** 1 + 8
- **On-chip Memories**
 - A level 2 Memory, shared among all cores
 - A level 1 Memory, shared by the 8-cores cluster
- **cluster-DMA:** A multi-channel 1D/2D DMA, controlling the transactions between the L2 and L1 memories
- **micro-DMA:** A smart, lightweight and completely autonomous DMA () capable of handling complex I/O scheme
- **Bus+Peripherals:** HyperBus, I2S, CPI, timers, SPI, GPIOs, etc...

NB: this is the architecture you find on our nano-drones and GAP boards!


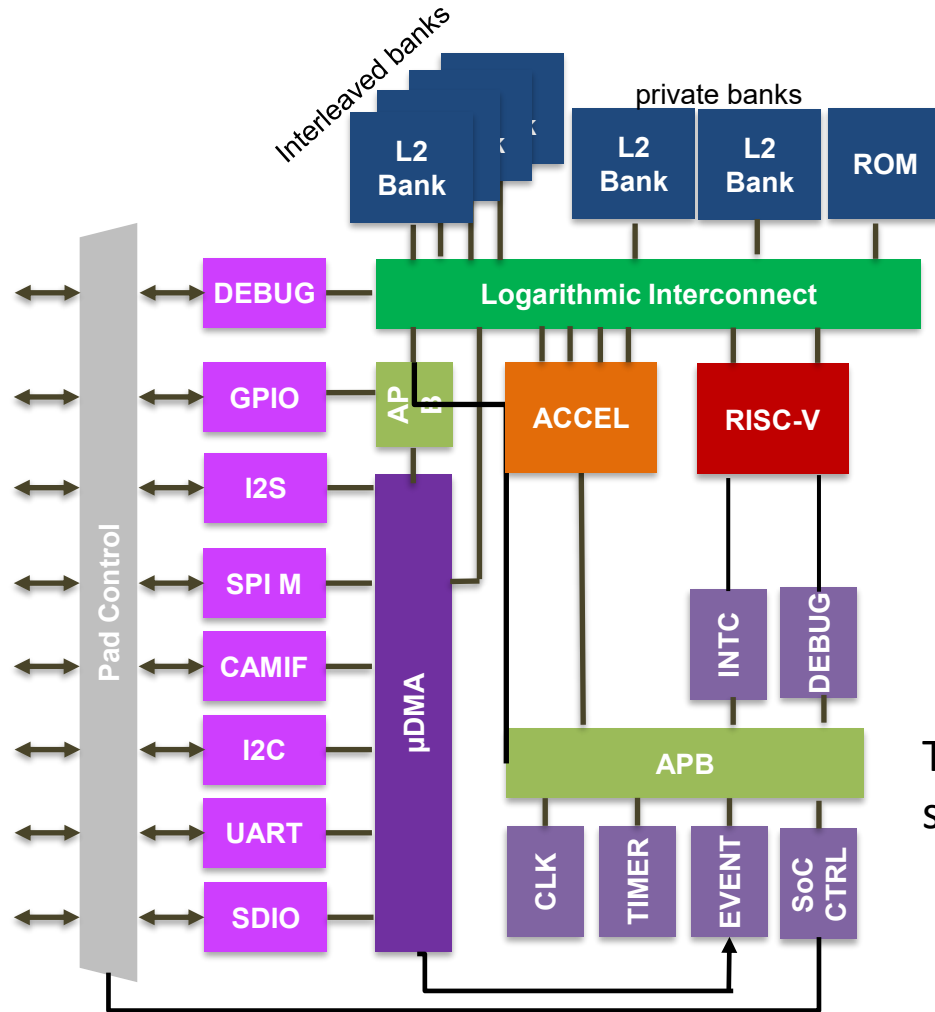
GitHub HW Project: <https://github.com/pulp-platform/pulp>
HW Documentation: <https://raw.githubusercontent.com/pulp-platform/pulp/master/doc/datasheet.pdf>



PULP: a Parallel Ultra-Low-Power computing platform


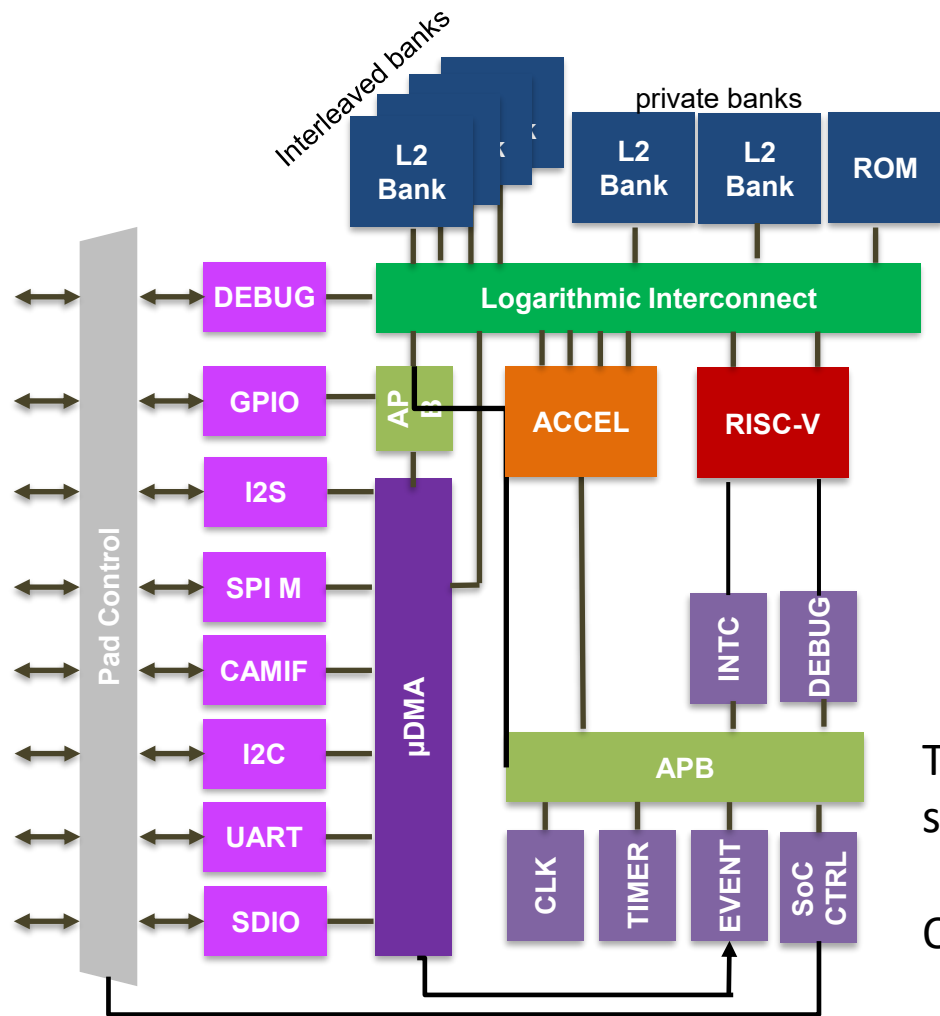


PULP: a Parallel Ultra-Low-Power computing platform



Target a **Shared-Memory** parallel programming model: 4—16 DSP cores sharing directly a **Shared L1 Memory** (*Tightly Coupled Data Mem* or *TCDM*)

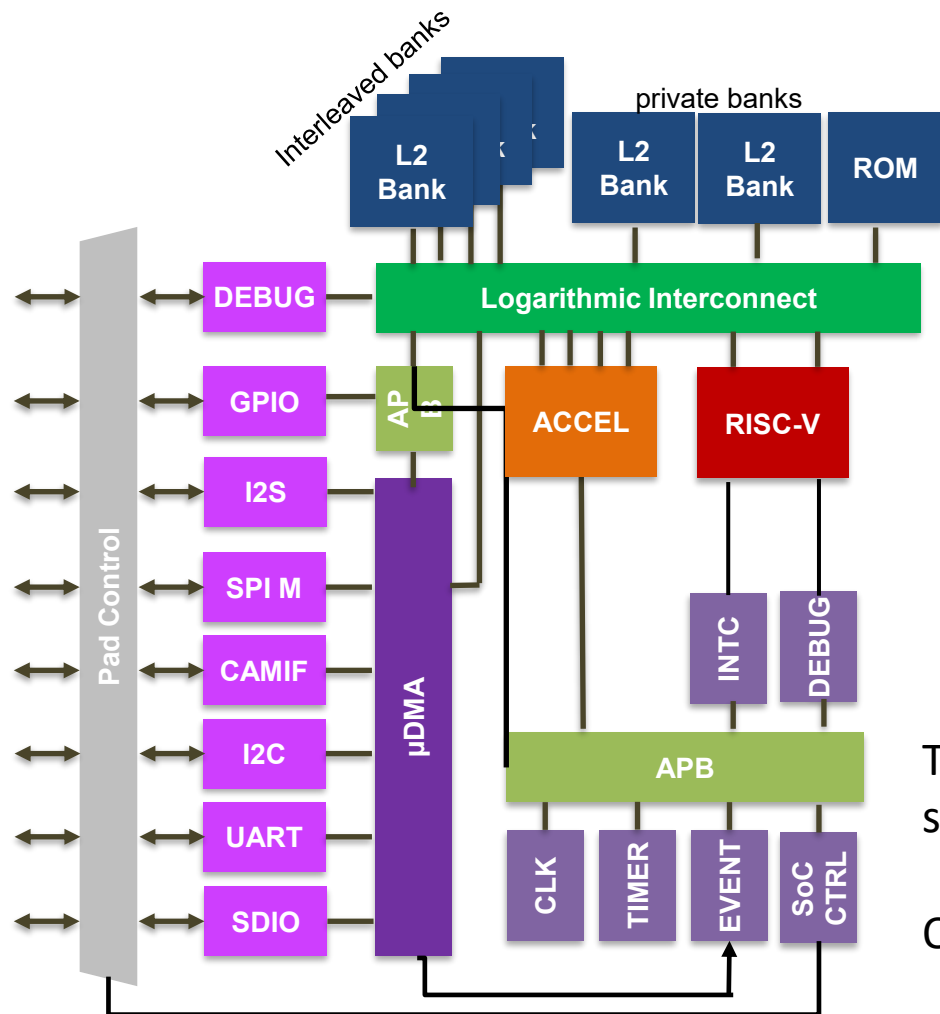
PULP: a Parallel Ultra-Low-Power computing platform



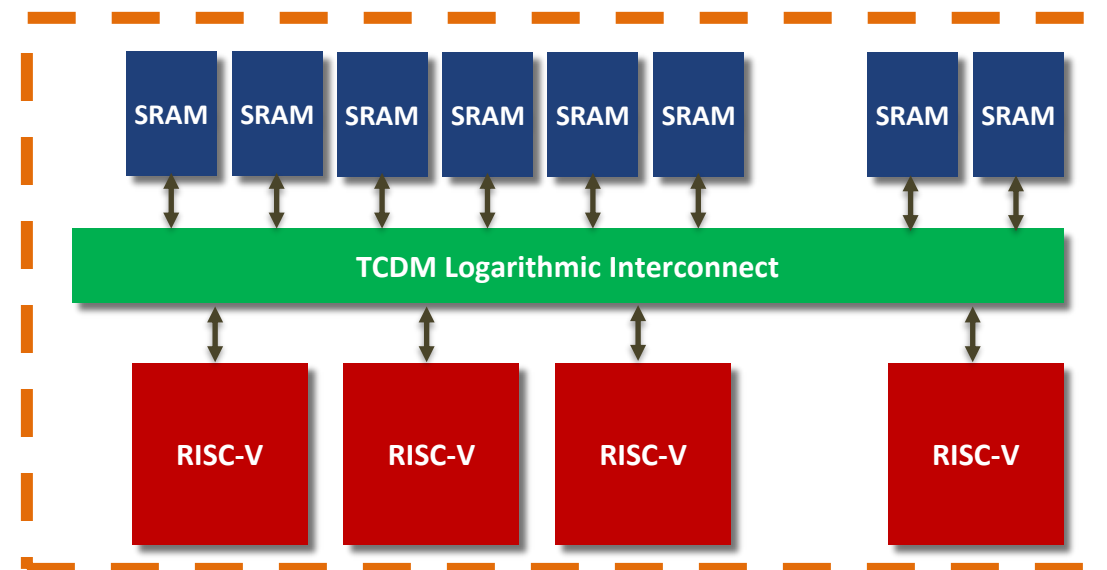
Target a **Shared-Memory** parallel programming model: 4—16 DSP cores sharing directly a **Shared L1 Memory** (*Tightly Coupled Data Mem* or *TCDM*)

Organize memory in **Multiple Banks** → concurrent access

PULP: a Parallel Ultra-Low-Power computing platform



PULP Cluster

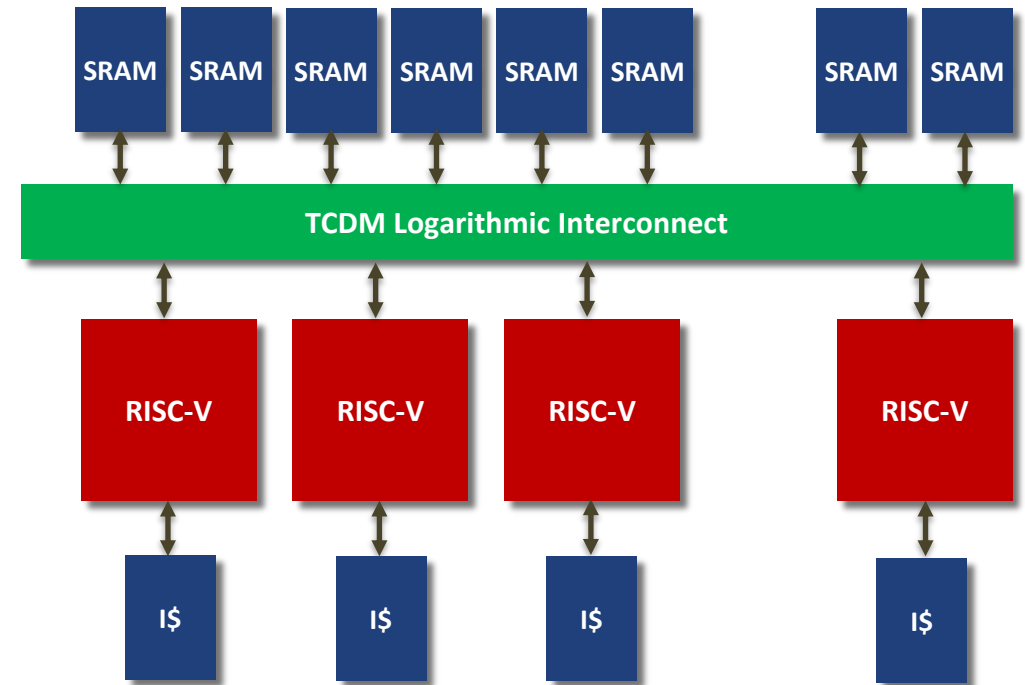



Target a **Shared-Memory** parallel programming model: 4—16 DSP cores sharing directly a **Shared L1 Memory** (*Tightly Coupled Data Mem* or *TCDM*)

Organize memory in **Multiple Banks** ? concurrent access

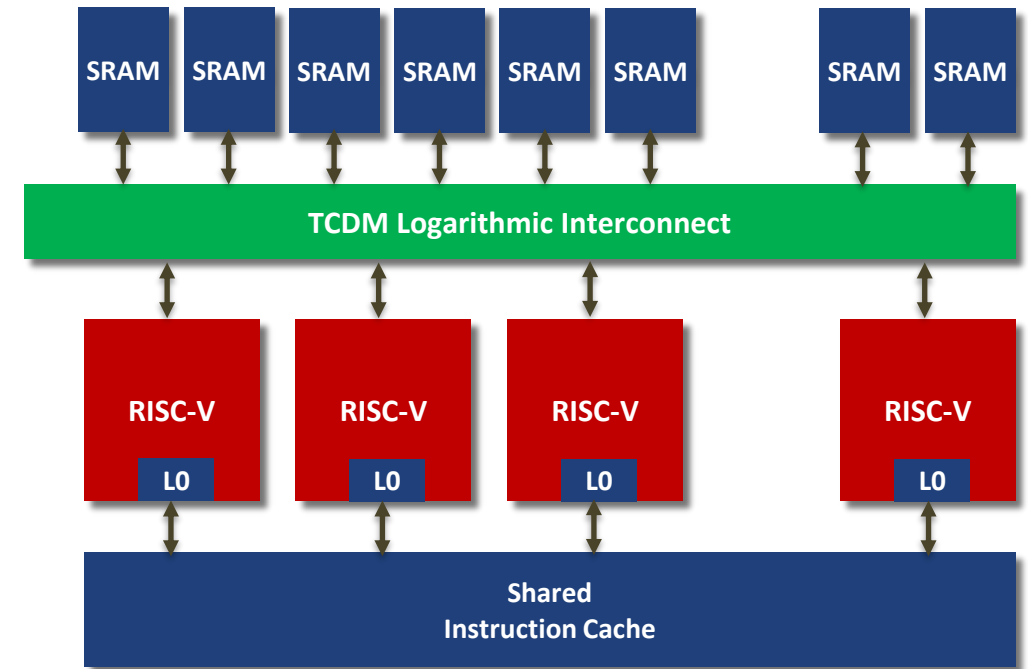
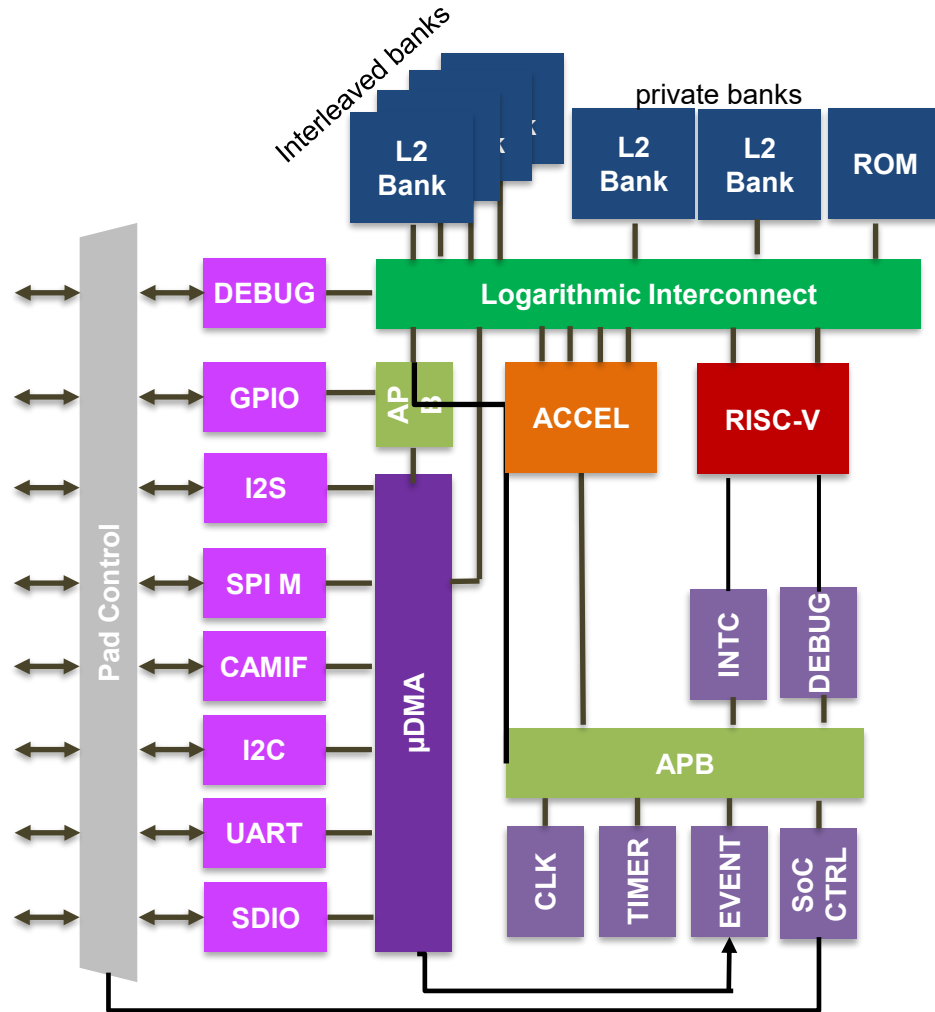
Full connectivity **1-cycle Access Crossbar**

PULP: a Parallel Ultra-Low-Power computing platform



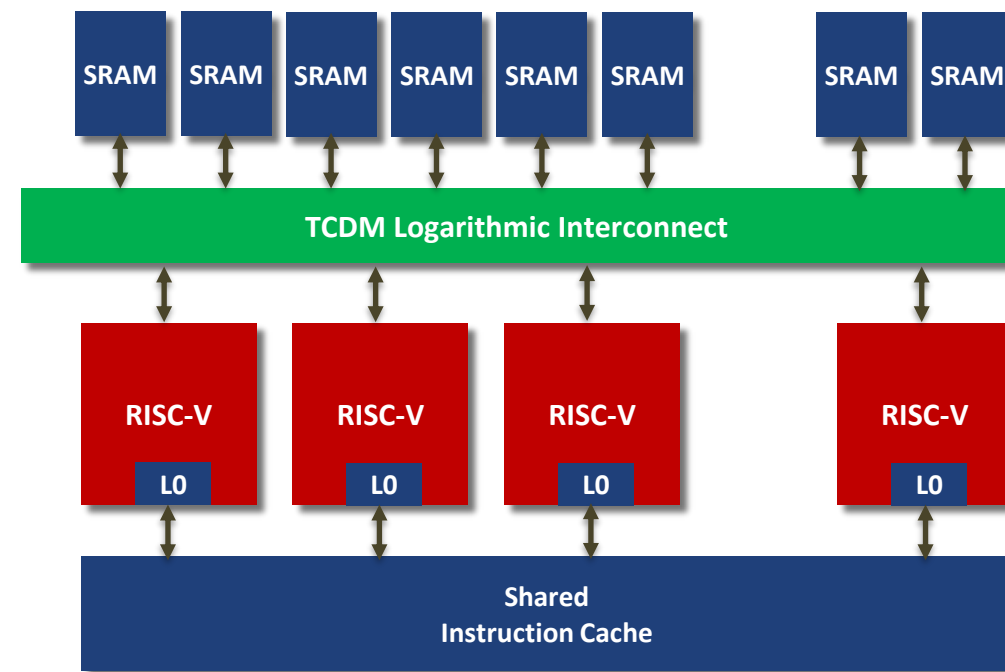
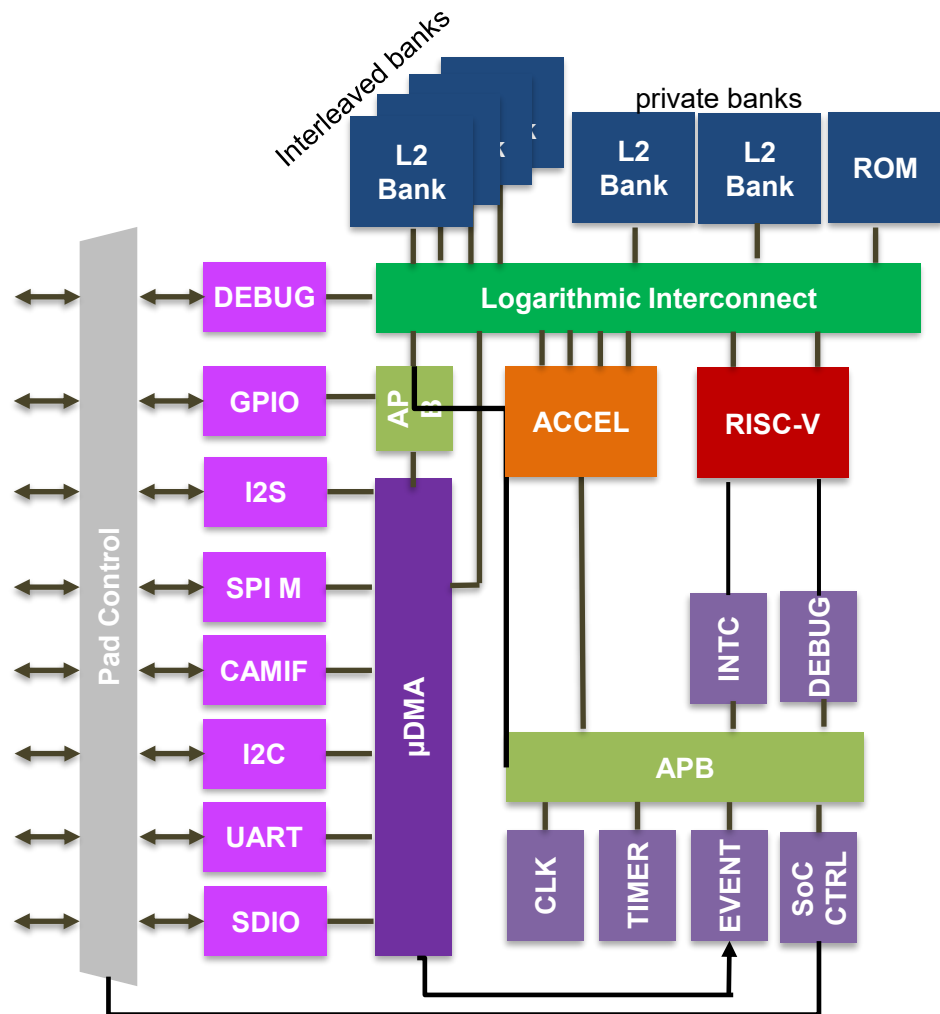
Multiple separate instruction streams (hardware threads):
MIMD execution scheme...

PULP: a Parallel Ultra-Low-Power computing platform



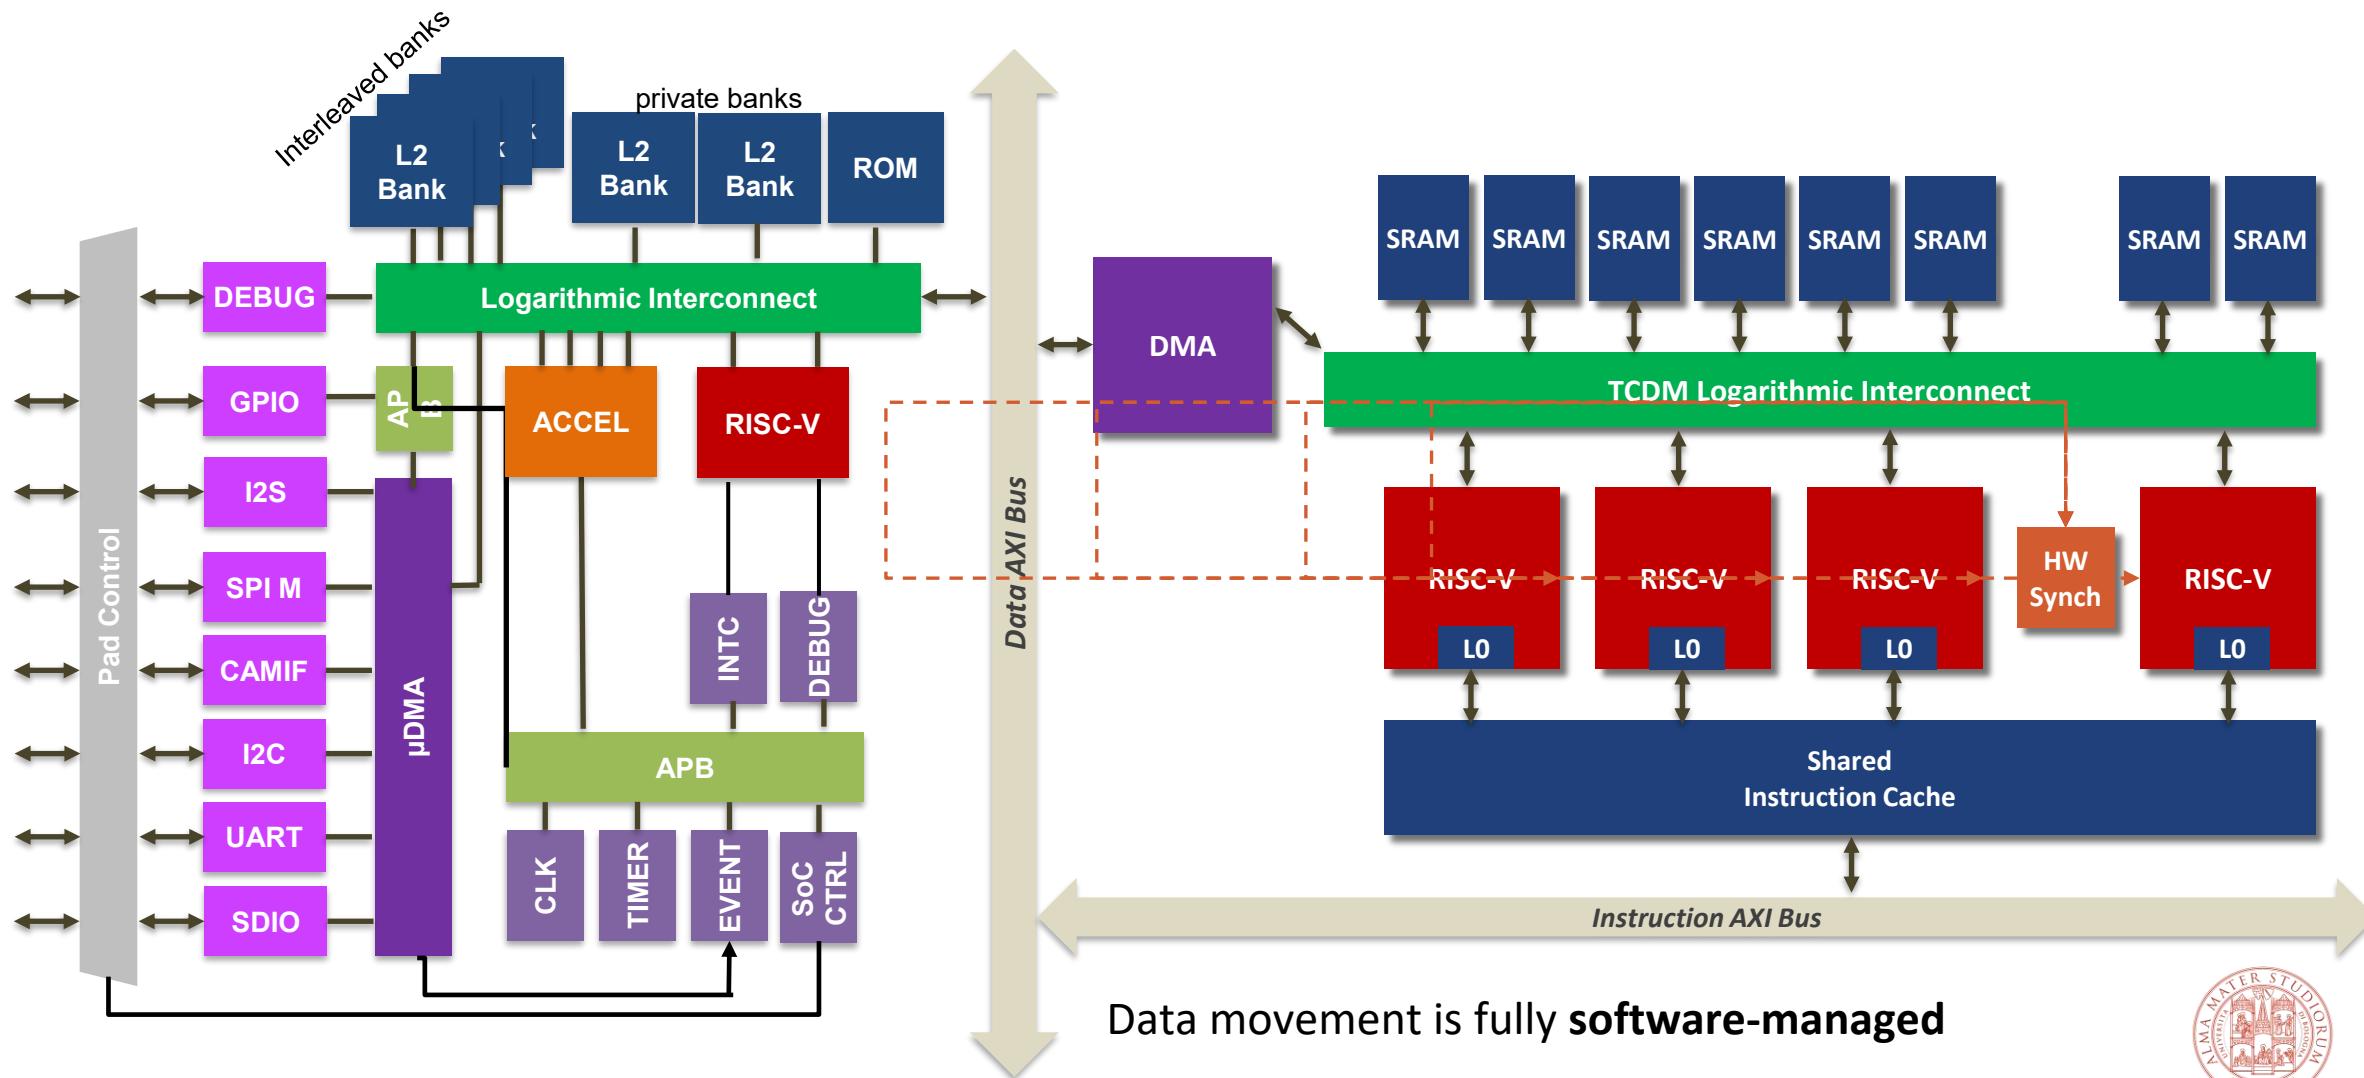
Multiple separate instruction streams (hardware threads):
MIMD execution scheme...
... but optimized for a single parallel program (**SPMD**)

PULP: a Parallel Ultra-Low-Power computing platform



SPMD executes multiple instances of the same program independently, where each program works on a different portion of the data

PULP: a Parallel Ultra-Low-Power computing platform



PMSIS: how to manage a device lifecycle

- Configuration and initialization (*device specific*)
conf_init()
open_from_conf()
- Prepare the device for usage:
open()
- Perform required operations (*device specific*)
- Release the resources
close()

PMSIS: Using the cluster as a device

You will see this in main.c

```
struct pi_device cluster_dev = {0};  
struct pi_cluster_conf conf;  
struct pi_cluster_task cluster_task = {0};
```

} PMSIS data structures

```
// task allocation  
pi_cluster_task(&cluster_task, cluster_entry, NULL);
```

} Create a task for the cluster

```
// init the cluster  
pi_cluster_conf_init(&conf);  
pi_open_from_conf(&cluster_dev, &conf);
```

 Function pointer

} Initialize the cluster device

```
// open the cluster  
if (pi_cluster_open(&cluster_dev)) return -1;
```

} Open the cluster device

```
// offload an entry point to the cluster  
pi_cluster_send_task_to_cl(&cluster_dev, &cluster_task);
```

} Execute code on the cluster

```
// releasing the cluster  
pi_cluster_close(&cluster_dev);
```

} Release the cluster device

Executing on the cluster: Fork/join + SPMD

```
static void cluster_entry(void *arg)
{
    printf("Hello from cluster\n");
    // ...

    pi_cl_team_fork(NUM_CORES, cluster_fn, (void *) &args);

    // ...
}
```

Data pointer

Function pointer

} Executed on core 0

} Fork the execution of the same function on
NUM_CORES cores

} Executed on core 0

SPMD executes multiple instances of the same program independently, where each program works on a different portion of the data

Synchronization: Barriers and Critical sections

- Barriers (*used for intermediate and final join points*):

```
pi_cl_team_barrier();
```

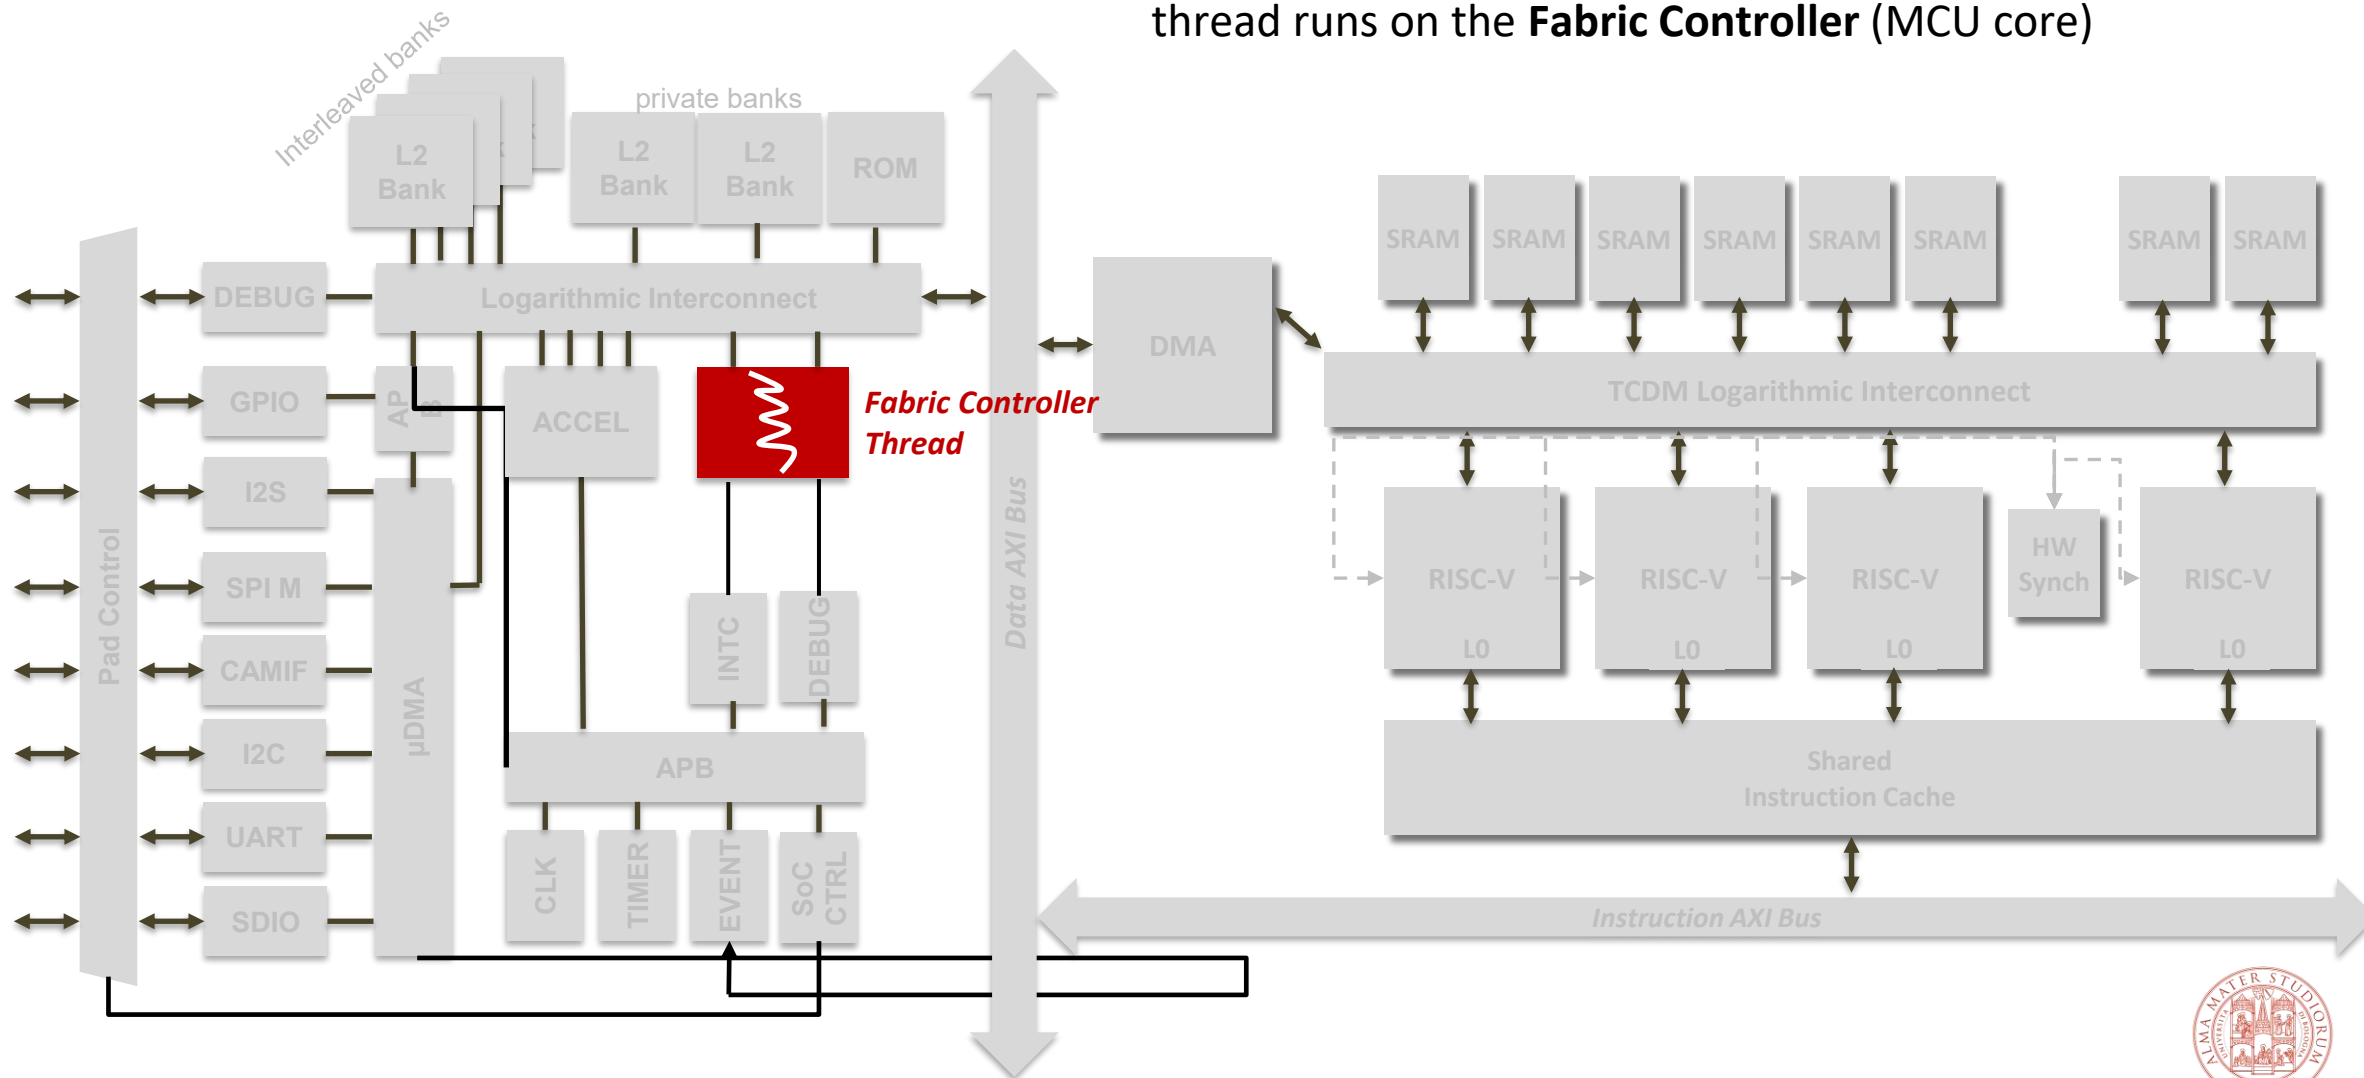
- Critical sections (*used to avoid critical races*):

```
pi_cl_team_critical_enter();  
// code in the critical section  
pi_cl_team_critical_exit();
```




PULP execution model

Cluster is **inactive** and **clock-gated** at boot: a single thread runs on the **Fabric Controller** (MCU core)



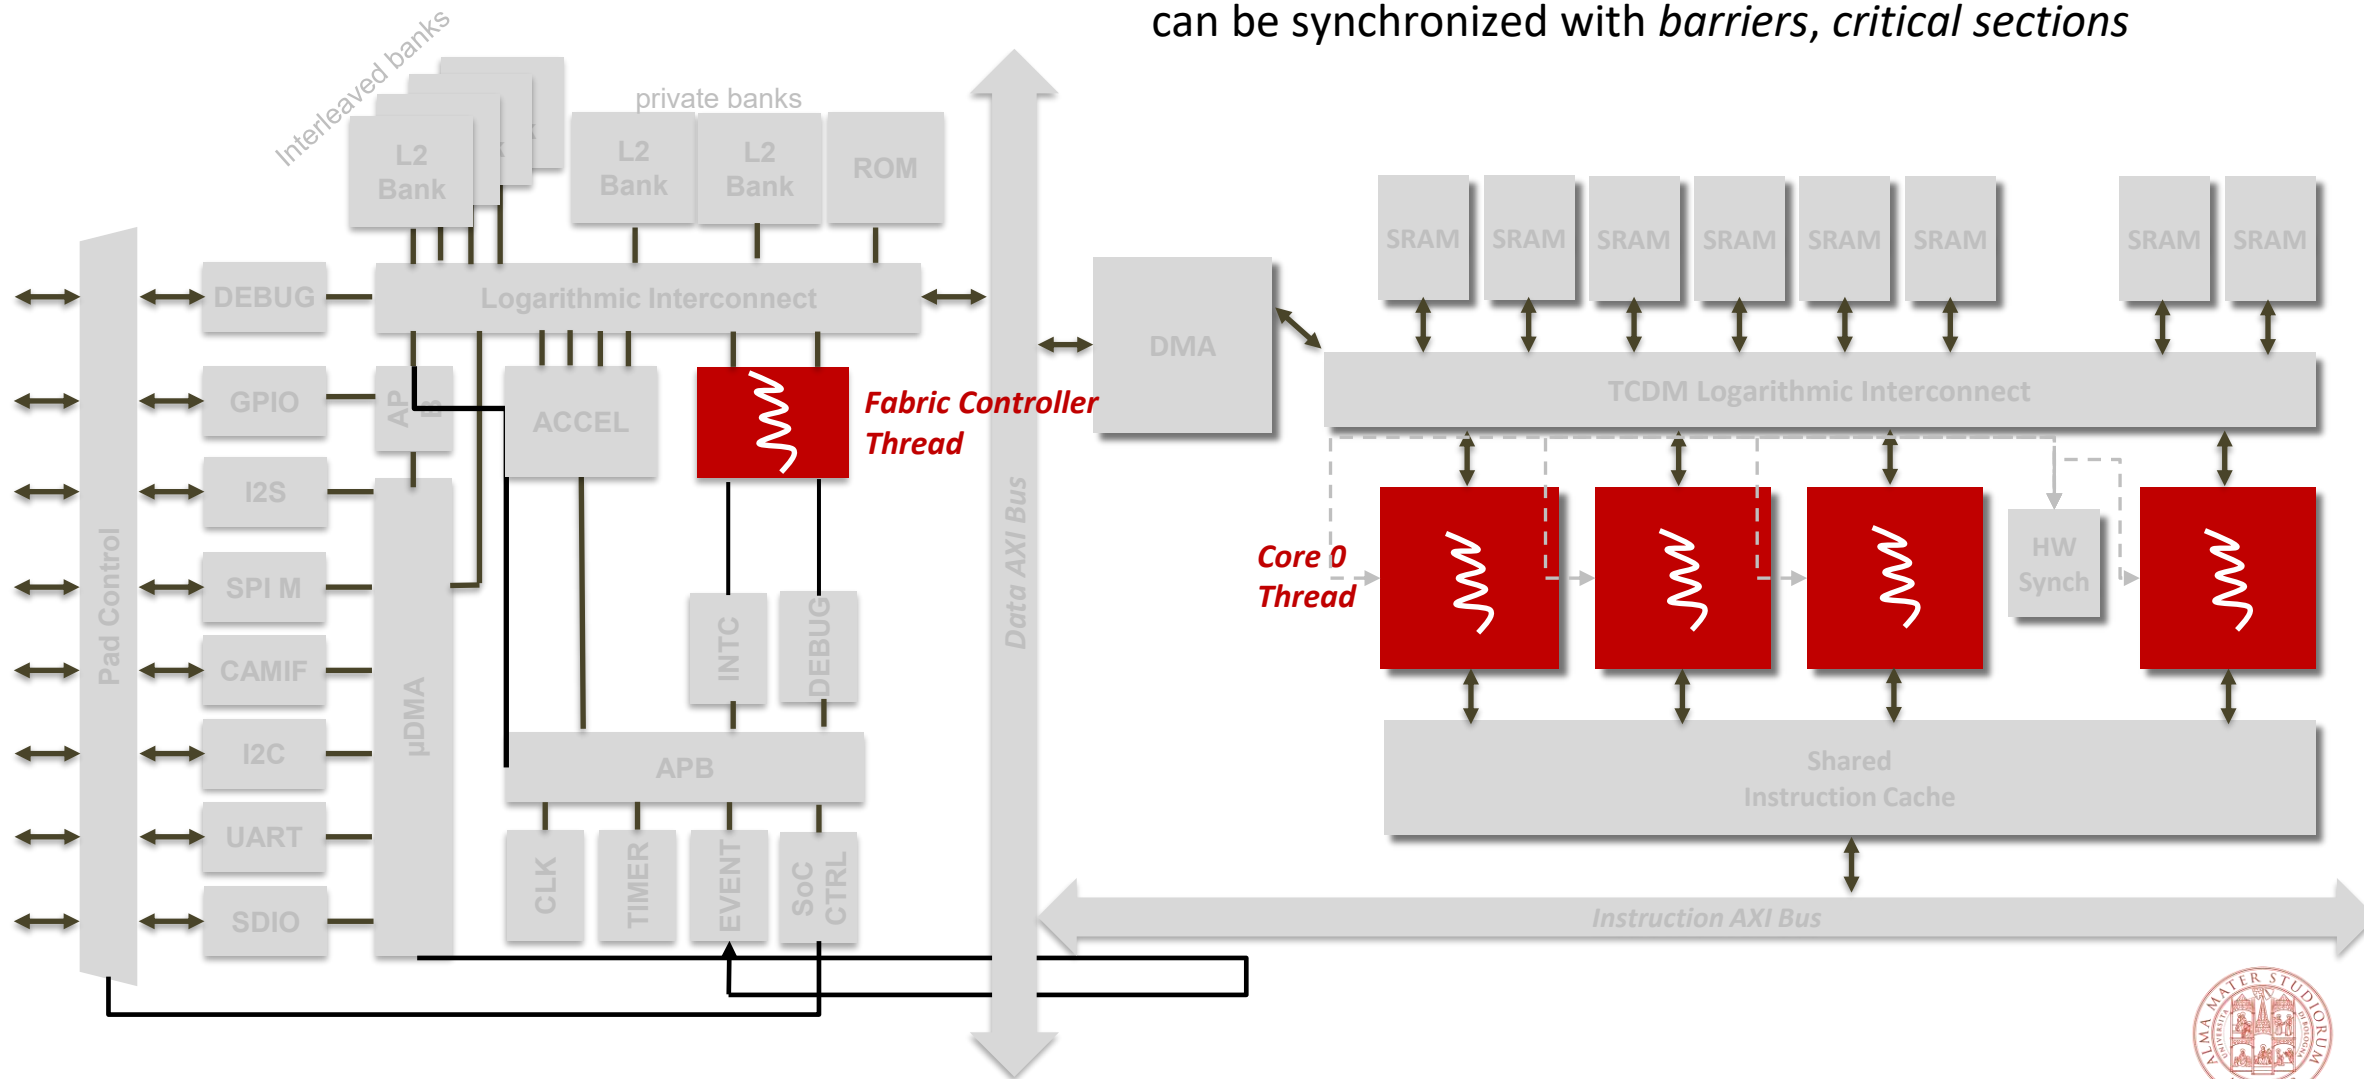
PULP execution model

Activate the cluster (clock it!) and **call** a function on the **cluster core 0**



PULP execution model

Fork an **execution team** on multiple cluster cores; they can be synchronized with *barriers*, *critical sections*






ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

TASK1: MatMul

FROM LAB02: Matrix-vector product




```
// generic matrix-vector multiplication
int gemv(int N, int M, float * mat, float *vec, float * output_vec){

    for (int i=0; i<N; i++){
        for (int j=0; j<M; j++){
            vec_o[i] += mat_i[i*size_M+j] * vec_i[j];
        }
    }

}
```

Today: Matrix Multiplication (MatMul)



```
// generic matrix multiplication
void gemm(int * MatA, int * MatB, int* MatC, int NN, int MM, int KK){
    // task to profile
    for (int i = 0; i < NN; i++) {
        for (int j = 0; j < MM; j++) {
            int acc = 0;
            for (int k = 0; k < KK; k++) {
                acc += MatA[i*KK+k] * MatB[k*MM+j];
            } //k
            MatC[i*MM+j] = acc;
        } //j
    } //i
```

MACs = N*K*M

Matrix Multiplication and Parallelization

Now try to execute the parallelized version of matrix multiplication:

```
$ cd matmul_parallelization/  
$ make clean all run CORES=<1 to 8>
```

- Follow the assignment document.



TASK1.1: Matrix Multiplication and Parallelization

```
// generic matrix multiplication
void gemm(int * MatA, int * MatB, int* MatC, int NN, int MM, int KK){
    uint32_t i, core_id, i_chunk, i_start, i_end;
```

```
    core_id = pi_core_id();
```

```
    i_chunk = (NN + NUM_CORES-1) / NUM_CORES;
```

```
    i_start = core_id * i_chunk;
```

```
    i_end    = i_start + i_chunk < NN ? i_start + i_chunk : NN;
```

} Divide the workload in adjacent blocks (*chunks*)
and compute their bounds [i_start, i_end]

```
// task to profile
```

```
for (i = i_start; i < i_end; i++) {
```

```
    for (int j = 0; j < MM; j++) {
```

```
        int acc = 0;
```

```
        for (int k = 0; k < KK; k++) {
```

```
            acc += MatA[i*KK+k] * MatB[k*MM+j];
```

```
        } //k
```

```
        MatC[i*MM+j] = acc;
```

```
    }//j
```

```
}//I
```

```
pi_cl_team_barrier();
```

```
}
```

We talked before about SPMD:

executes multiple instances of the same
program independently, where each program
works on a different portion of the data

TASK1.1: Matrix Multiplication and Parallelization

```
// generic matrix multiplication
void gemm(int * MatA, int * MatB, int* MatC, int NN, int MM, int KK){
    uint32_t i, core_id, i_chunk, i_start, i_end;
```

```
    core_id = pi_core_id();
```

```
    i_chunk = (NN + NUM_CORES-1) / NUM_CORES;
```

```
    i_start = core_id * i_chunk;
```

```
    i_end   = i_start + i_chunk < NN ? i_start + i_chunk : NN;
```

} Divide the workload in adjacent blocks (*chunks*)
and compute their bounds [i_start, i_end]

```
// task to profile
```

```
for (i = i_start; i < i_end; i++) {
```

```
    for (int j = 0; j < MM; j++) {
```

```
        int acc = 0;
```

```
        for (int k = 0; k < KK; k++) {
```

```
            acc += MatA[i*KK+k] * MatB[k*MM+j];
```

```
        } //k
```

```
        MatC[i*MM+j] = acc;
```

```
    } //j
```

```
} //I
```

```
pi_cl_team_barrier();
```

```
}
```

We talked before about SPMD:

executes multiple instances of the same program independently, where each program works on a different portion of the data

core_id is used to divide portion of the data among all cores


TASK1.2: different input sizes

- We parallelize on the number of rows N .
- We have 8 cores

If the number of rows is 4: \rightarrow only 4 cores busy



If the number of rows is 8 \rightarrow 8 cores busy



Task 1.3: Matrix Multiplication and Parallelization w/ Manual Unrolling

Manual (or static) **loop unrolling** helps reducing the total number of instructions!.

```
void gemm_unroll_1x4(int * MatA, int * MatB, int* MatC, int
NN, int MM, int KK){
```

```
    uint32_t core_id, i_chunk, i_start, i_end;
    uint32_t i = 0;
```

```
    core_id = pi_core_id();
    i_chunk = (NN + NUM_CORES-1) / NUM_CORES;
    i_start = core_id * i_chunk;
    i_end    = i_start + i_chunk < NN ? i_start + i_chunk :
NN;
```


```
        < unrolled loop here >
```

```
    }
```

```
...
for (i = i_start; i < i_end; i++) {
    for (int j = 0; j < MM; j=j+4) {
        int acc0 = 0;
        int acc1 = 0;
        int acc2 = 0;
        int acc3 = 0;
        for (int k = 0; k < KK; k++) {
            int shared_op = MatA[i*KK+k];
            int idx = k*MM+j;
            acc0 += shared_op * MatB[idx];
            acc1 += shared_op * MatB[idx+1];
            acc2 += shared_op * MatB[idx+2];
            acc3 += shared_op * MatB[idx+3];
        } //k
        MatC[i*MM+j] = acc0;
        MatC[i*MM+j+1] = acc1;
        MatC[i*MM+j+2] = acc2;
        MatC[i*MM+j+3] = acc3;
    } //j
} //i
```

Since you unroll 4 operations,
you should cycle 4 by 4.

Matrix Multiplication: Baseline



Internal loop

4096 (16x16x16) iterations


4 instructions \approx **16834** ops

```
lp.setup      x1,a6,1c0089e4

lbu           t4,0(t1)    // load in2
p.lbu         t5,1(t3!)   // load in1
add           t1,t1,a4     // in2 incr
p.mac         a7,t5,t4     // out += in2*in1
```

2 LOAD + 1 ADD + 1 MAC

```
for (int i = 0; i < MN; i++) {
    for (int j = 0; j < NN; j++) {
        int acc = 0;
        for (int k = 0; k < KN; k++) {
            acc += in1[i*KN+k] * in2[k*NN+j];
        } //k
        out3[i*NN+j] = acc;
    } //j
} //i
```



Number of Instructions: **18867**

Clock Cycles: **18911**

CPI: **1.00**

MAC/cyc: **0.21**


$MN = NN = KN = 16$

$MAC = 16^3 = 4096$

Matrix Multiplication: Unrolling I

Loop Unrolling to enable data reuse and reduce number of load instructions!

```
for (int i = 0; i < MN; i++) {  
    for (int j = 0; j < NN; j=j+2) {  
        int acc0 = 0;  
        int acc1 = 0;  
        for (int k = 0; k < KN; k++) {  
            char shared_op = in1[i*KN+k];  
            acc0 += shared_op * in2[k*NN+j];  
            acc1 += shared_op * in2[k*NN+j+1];  
        } //k  
        out3[i*NN+j] = acc0;  
        out3[i*NN+j+1] = acc1;  
    } // loop j  
} // loop i
```



Internal loop

2048 (16x8x16) iterations

7 instructions \approx **14336** ops

```
p.lbu      t3, 1(t6!)  
lb         t2, 0(t1)  
lb         t0, 0(a7)  
add        t1, t1, a4  
p.mac      t4, t3, t2  
add        a7, a7, a4  
p.mac      t5, t3, t0
```

3 LOAD + 2 ADD + 2 MAC

Number of Instructions: **16282**

Clock Cycles: **16326**

CPI: **1.00**

MAC/cyc: **0.25**


1.15x!

Matrix Multiplication: Unrolling II

```

for (int i = 0; i < MN; i++) {
    for (int j = 0; j < NN; j=j+4) {
        int acc0 = 0;
        int acc1 = 0;
        int acc2 = 0;
        int acc3 = 0;
        for (int k = 0; k < KN; k++) {
            char shared_op = in1[i*KN+k];
            acc0 += shared_op * in2[k*NN+j];
            acc1 += shared_op * in2[k*NN+j+1];
            acc2 += shared_op * in2[k*NN+j+2];
            acc3 += shared_op * in2[k*NN+j+3];
        } //k
        out3[i*NN+j]    = acc0;
        out3[i*NN+j+1] = acc1;
        out3[i*NN+j+2] = acc2;
        out3[i*NN+j+3] = acc3;
    } //j
} //i

```



Internal loop

1024 (16x4x16) iterations

10 instructions \approx **10240** ops

```

p.lbu    a7,1(t0!)
lb       s0,0(a6)
lb       a0,1(a6)
lb       a2,2(a6)
lb       t2,3(a6)
p.mac    t3,a7,s0
add      a6,a6,a4
p.mac    t4,a7,a0
p.mac    t5,a7,a2
p.mac    t6,a7,t2

```

5 LOAD + 1 ADD + 4MAC

Number of Instructions: **11579**

Clock Cycles: **11719**

CPI: **1.01**

MAC/cyc: **0.35**

1.4x!

Matrix Multiplication: Unrolling III

```


for (int i = 0; i < MN; i=i+2) {
    for (int j = 0; j < NN; j=j+4) {
        acc0 = acc1 = acc2 = acc3 = 0;
        acc4 = acc5 = acc6 = acc7 = 0;

        for (int k = 0; k < KN; k++) {
            char shared_op = in1[i*KN+k];
            char a = in2[k*NN+j];
            char b = in2[k*NN+j+1];
            char c = in2[k*NN+j+2];
            char d = in2[k*NN+j+3];

            acc0 += shared_op * a;
            acc1 += shared_op * b;
            acc2 += shared_op * c;
            acc3 += shared_op * d;

            shared_op = in1[(i+1)*KN+k];
            acc4 += shared_op * a;
            acc5 += shared_op * b;
            acc6 += shared_op * c;
            acc7 += shared_op * d;
        } //k
        out3[i*NN+j] = acc0;      out3[i*NN+j+1] = acc1;
        out3[i*NN+j+2] = acc2;    out3[i*NN+j+3] = acc3;
        out3[(i+1)*NN+j] = acc4;  out3[(i+1)*NN+j+1] = acc5;
        out3[(i+1)*NN+j+2] = acc6; out3[(i+1)*NN+j+3] = acc7;
    } //j
} //I

```



Internal loop

512 (8x4x16) iterations

15 instructions ≈ 7680 ops

p.lbu	a2,1(s2!)	p.mac	t4,a2,a7
p.lbu	a3,1(s3!)	p.mac	t5,a2,a6
lb	t1,0(a5)	p.mac	t6,a2,a0
lb	a7,1(a5)	p.mac	t0,t1,a3
lb	a6,2(a5)	p.mac	t2,a7,a3
lb	a0,3(a5)	p.mac	s0,a6,a3
p.mac	t3,a2,t1	p.mac	s1,a0,a3
add	a5,a5,a4		

6 LOAD + 1 ADD + 8 MAC

Number of Instructions: **8768**
 Clock Cycles: **8812**
 CPI: **1.01**
 MAC/cyc: **0.464**

1.33x!

Matrix Multiplication: Unrolling IV

```

for (int i = 0; i < MN; i=i+4) {
    for (int j = 0; j < NN; j=j+4) {
        ...
        for (int k = 0; k < KN; k++) {
            char shared_op = in1[i*KN+k];
            char a = in2[k*NN+j];
            char b = in2[k*NN+j+1];
            char c = in2[k*NN+j+2];
            char d = in2[k*NN+j+3];

            acc0 += shared_op * a;
            acc1 += shared_op * b;
            acc2 += shared_op * c;
            acc3 += shared_op * d;

            shared_op = in1[(i+1)*KN+k];
            acc4 += shared_op * a;
            acc5 += shared_op * b;
            acc6 += shared_op * c;
            acc7 += shared_op * d;

            shared_op = in1[(i+2)*KN+k];
            acc8 += shared_op * a;
            acc9 += shared_op * b;
            acc10 += shared_op * c;
            acc11 += shared_op * d;

            shared_op = in1[(i+3)*KN+k];
            acc12 += shared_op * a;
            acc13 += shared_op * b;
            acc14 += shared_op * c;
            acc15 += shared_op * d;
        }
    }
}


```

Number of Instructions: **10546**
 Clock Cycles: **11342**
 CPI: **1.07**
 MAC/cyc: **0.361**

```

...
p.lbu    a7,1(s9!)
p.lbu    a6,1(s10!)
p.lbu    a4,1(s11!)
sw        t4,8(sp)
lw        t4,40(sp)
p.mac     s1,a2,a7
add        a5,a5,t4
lw        t4,20(sp)
p.mac     t2,a0,a7
p.mac     s0,a1,a7
p.mac     s2,a3,a7
lw        a7,16(sp)
p.mac     t4,a0,a4
p.mac     a7,a3,a6
sw        t4,20(sp)
p.mac     t6,t1,a2
...

```




The register file is not infinite! ☹️

For wide loop body, the compiler may cannot find **enough registers to map all the required variables**. If this happens, variables start getting stored on the *stack (spilling)*

- Performance drop!



Recap: Loop Unrolling of MatMult



Note: assumption of MN/NN multiple of $2/4$. *What if this was not the case?*

Disassembly code for unrolled matmul: no stalls

1. Parallel MatMul w/ manual unrolling.

- Replace *gemm* with *gemm_unroll*
- Plot #instr and #clk by varying CORES
- You can see it in the disassembly

Results with Manual Unroll:

5 LD + 4 MAC + 1 ADD

7741213546: 506585: [[34m/sys/board/chip/cluster/pe0/insn	[0m] gemm_unroll_1x4:71	M 1c008a42	p.mac	t3, a6, a2	t3=00000004 t3:00000002 a6:00000002 a2:00000001
7741233572: 506586: [[34m/sys/board/chip/cluster/pe0/insn	[0m] gemm_unroll_1x4:71	M 1c008a46	c.add	a3, a3, s0	a3=1000009c a3:1000005c s0:00000040
7741253598: 506587: [[34m/sys/board/chip/cluster/pe0/insn	[0m] gemm_unroll_1x4:70	M 1c008a48	p.mac	t1, a6, s9	t1=00000004 t1:00000002 a6:00000002 s9:00000001
7741273624: 506588: [[34m/sys/board/chip/cluster/pe0/insn	[0m] gemm_unroll_1x4:72	M 1c008a4c	p.mac	t4, a6, t2	t4=00000004 t4:00000002 a6:00000002 t2:00000001
7741293650: 506589: [[34m/sys/board/chip/cluster/pe0/insn	[0m] gemm_unroll_1x4:73	M 1c008a50	p.mac	t5, a6, t0	t5=00000004 t5:00000002 a6:00000002 t0:00000001
7741313676: 506590: [[34m/sys/board/chip/cluster/pe0/insn	[0m] gemm_unroll_1x4:68	M 1c008a30	p.lw	a6, 4(t6!)	a6=00000002 t6=10001628 t6:10001624 PA:10001624
7741333702: 506591: [[34m/sys/board/chip/cluster/pe0/insn	[0m] gemm_unroll_1x4:71	M 1c008a34	c.lw	a2, 4(a3)	a2=00000001 a3:1000009c PA:100000a0
7741393780: 506594: [[34m/sys/board/chip/cluster/pe0/insn	[0m] gemm_unroll_1x4:72	M 1c008a36	lw	t2, 8(a3)	t2=00000001 a3:1000009c PA:100000a4
7741433832: 506596: [[34m/sys/board/chip/cluster/pe0/insn	[0m] gemm_unroll_1x4:73	M 1c008a3a	lw	t0, 12(a3)	t0=00000001 a3:1000009c PA:100000a8
7741453858: 506597: [[34m/sys/board/chip/cluster/pe0/insn	[0m] gemm_unroll_1x4:70	M 1c008a3e	lw	s9, 0(a3)	s9=00000001 a3:1000009c PA:1000009c

Write parallel code

Try on your own!

1. Parallel MatMul w/ manual unrolling.

- Create a new folder (`cp -r matmul_parallelization/ matmul_unroll/`)
- Replace `gemm` with `gemm_unroll_1x4`
- Adjust `gemm_unroll_1x4` to combine loop unrolling with parallelization
- Plot `#instr` and `#clk` by varying CORES
- Analyze your measurements against disassembly (do the measured number of `#instructions` and `#clk` make sense?)

Results with Manual Unroll:

Number of Instructions: 33097

Clock Cycles: 33956 | **1 core** execution

Number of Instructions: 4190

Clock Cycles: 4548 | **8 cores** execution

S.U. = 7.90

Use traces to verify the amount of INSTR and CYC:

```
$ make clean all run runner_args='--trace=cluster/pe0/insn:trace.txt'
```

If you perform a `ctrl+f` in the **trace file**, searching for your **gemm** function name, you will see a corresponding number of instructions!



Reference on Makefile rules

The Makefile is a “recipe” used to call the compiler/linker and in this case also to run the program. You can also chain several rules (e.g., **make clean all run**) and pass options (e.g., **runner_args="--vcd"**)

Remove previous build

```
make clean
```

Build program (calling compiler + linker)

```
make all
```

Run the program

```
make run
```

Run options: you can change them by adding **runner_args="OPTIONS"**

```
make run runner_args="--vcd --event=.*" # visual trace in GTKwave
```

```
make run runner_args="--trace=.*insn.*" > trace.log # written trace of instructions
```

To open GTKwave, run the line that is visible in the log (**gtkwave some_long_string.gtkwave**) in the terminal

Disassembling:

```
make dis > test.S # disassemble without inlined source code
```

```
/pulp/pkg/pulp_riscv_gcc/bin/riscv32-unknown-elf-objdump -D -S > test.S # disassemble with inlined source code
```





ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

TASK2: FC

DNN Fully-Connected Layer on PULP

```
$ cd fully_connected/  
$ make clean all run CORES=<1 to 8>
```



$$Y = W \cdot X \rightarrow y[i] = \sum_{k=0}^M w[i, k] \cdot x[k]$$

matrix-vector multiplication

MAC operation!



$int32$ $int8$

- How many MACs per layer?
- How many cycles on a single core?
- How many cycles on a multi core platform?
 - Compute the speedup
- How many instructions? Check against the assembly code!

```
$ make dis
```

Benchmarking the Fully-Connected Layer

Number of MAC =
M x N =
CH_IM_IN x
CH_IM_OUT =
16384



# cores	#clk	#clk/ MAC	#instr	# instr/ MAC
1				
2				
4				
8				

```
1c008bd8 <pulp_nn_linear_u8_i32_i8>:
.....
1c008c2e: 0067c0fb lp.setup x1,a5,1c008c3a <pulp_nn_linear_u8_i32_i8+0x62>
1c008c32: 00134e8b p.lbu t4,1(t1!)
1c008c36: 00188e0b p.lb t3,1(a7!)
1c008c3a: 43ce8833 p.mac a6,t4,t3
.....
```


Check also the traces!

```
$ make clean all run CORES=8 runner_args="-trace=cluster/pe0/insn:log.txt"
```

TASK2.1: Vectorial Operations

Apply data-parallel processing on Vectorial Data

- **SIMD**: Single Instruction Multiple Data




The content of a 32-bit register can be interpreted as a vector of **2x16 bit values** or **4x8 bit values**

Remember LAB02: we quantized the network to 8 bits.
--> This enables SIMD !

Example: Vectorial Add

Vectors are packed in the same integer register-file!

- The instructions encode how to interpret the content of the register



Vectorial

Instructions of the
Xpulp extension

add rD, rs1, rs2	$rD = 0x03020100 + 0x0D0C0B0A$
pv.add.h rD, rs1, rs2	$rD[0] = 0x0100 + 0x0B0A$ $rD[1] = 0x0302 + 0x0D0C$
pv.add.b rD, rs1, rs2	$rD[0] = 0x00 + 0x0A$ $rD[1] = 0x01 + 0x0B$ $rD[2] = 0x02 + 0x0C$ $rD[3] = 0x03 + 0x0D$

2x 16-bit ADD in
one clock-cycles

4x 8-bit ADD in
one clock-cycles

Leftovers (vectorization of fully connected)

Replace this within the inner loop


```
v4u vecA;
v4s vecB;
// compute the vectorized dot products
for (int j=0; j<(dim_vec >> 2); j++) {
    vecA = *((v4u*)pA);
    vecB = *((v4s*)pB);
    sum = SumDotp4(vecA, vecB, sum);
    pA+=4;
    pB+=4;
}

// left over: handling the remaining input features
uint16_t col_cnt = dim_vec & 0x3;
while (col_cnt) {
    uint8_t inA = *pA;
    pA++;
    int8_t inB = *pB;
    pB++;
    sum += inA * inB;
    col_cnt--;
}
```

What if `dim_vec` is not
divisible by 4? ?Leftover

Handling the **left-over**
(`dim_vec %4 != 0`)

Benchmarking the Vectorized Fully-Connected Layer



Combining parallelization and vectorization:

- ~8x on 8 cores
- ~4x thanks to vectorization (4x less iteration!)

With SIMD

# cores	#clk	#MAC /clk	#instr	# instr/MAC
1				
2				
4				
8				

```

1c008bd8 <pulp_nn_linear_u8_i32_i8>:
.....
1c008c40: 0068c0fb    lp.setup    x1,a7,1c008c4c <pulp_nn_linear_u8_i32_i8+0x74>
1c008c44: 004ea80b    p.lw       a6,4(t4!)
1c008c48: 004e278b    p.lw       a5,4(t3!)
1c008c4c: a8f81357    pv.sdotusp.b t1,a6,a5
.....

```

\$ make dis

Similar to not-vectorized scheme:

- 3 instruction but 4 MAC now!



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA


TASK3: Conv2D

Convolution Operation

Convolution: Basic Processing Kernel of Convolutional Neural Networks

$$\mathbf{y} = \mathbf{w} * \mathbf{x}$$

$$\mathbf{y}[i, j] = \sum_{u_i=0}^{F_i-1} \sum_{u_j=0}^{F_j-1} \mathbf{w}[u_i, u_j] * \mathbf{x}[i + u_i, j + u_j] \quad (*)$$



(*) for math purists: “convolution” filters are technically **cross-correlations** (convolutions with flipped weights in both dimensions) in most cases of interest for digital signal processing

Convolution Operation

Convolution: Basic Processing Kernel of Convolutional Neural Network

$$y[i,j] = \sum_{u_i=0}^{F_i-1} \sum_{u_j=0}^{F_j-1} w[u_i, u_j] * x[i + u_i, j + u_j]$$





Image can be seen as Matrix



7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4

Kernel/Filter

*

1	0	-1
1	0	-1
1	0	-1

=

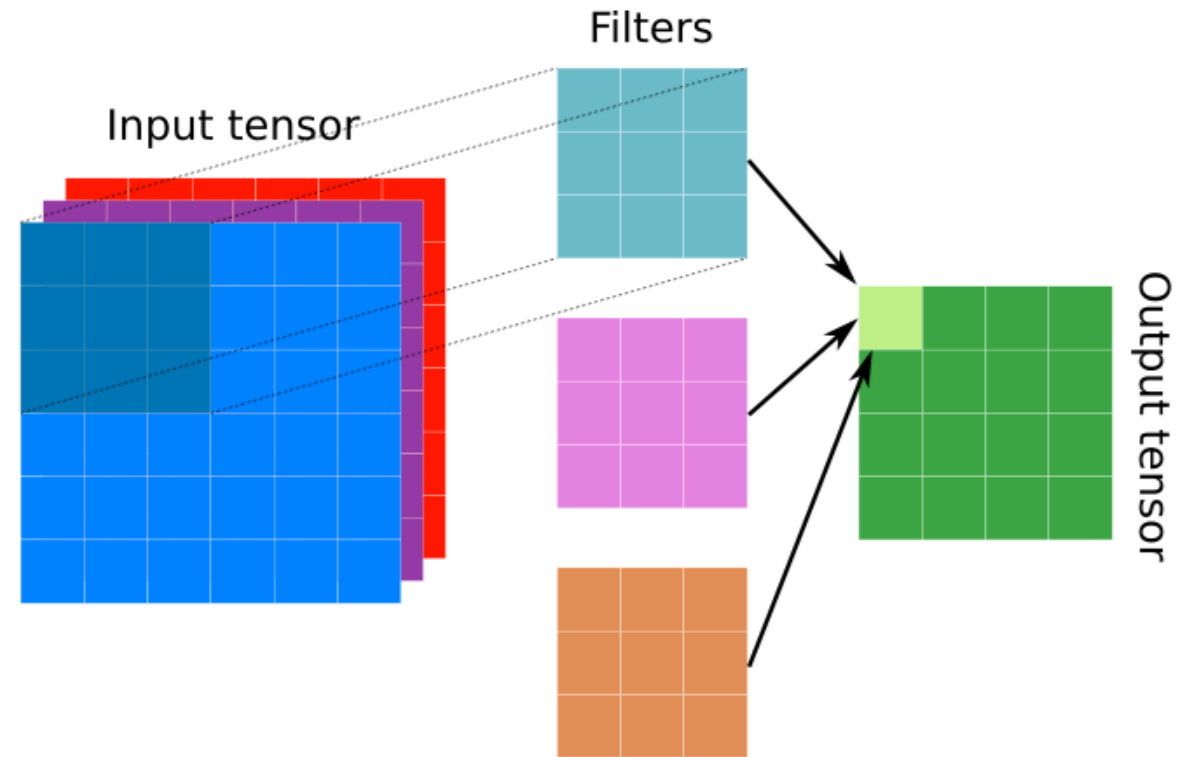
Feature Map

6		

$$7 \times 1 + 4 \times 1 + 3 \times 1 + 2 \times 0 + 5 \times 0 + 3 \times 0 + 3 \times -1 + 3 \times -1 + 2 \times -1 = 6$$

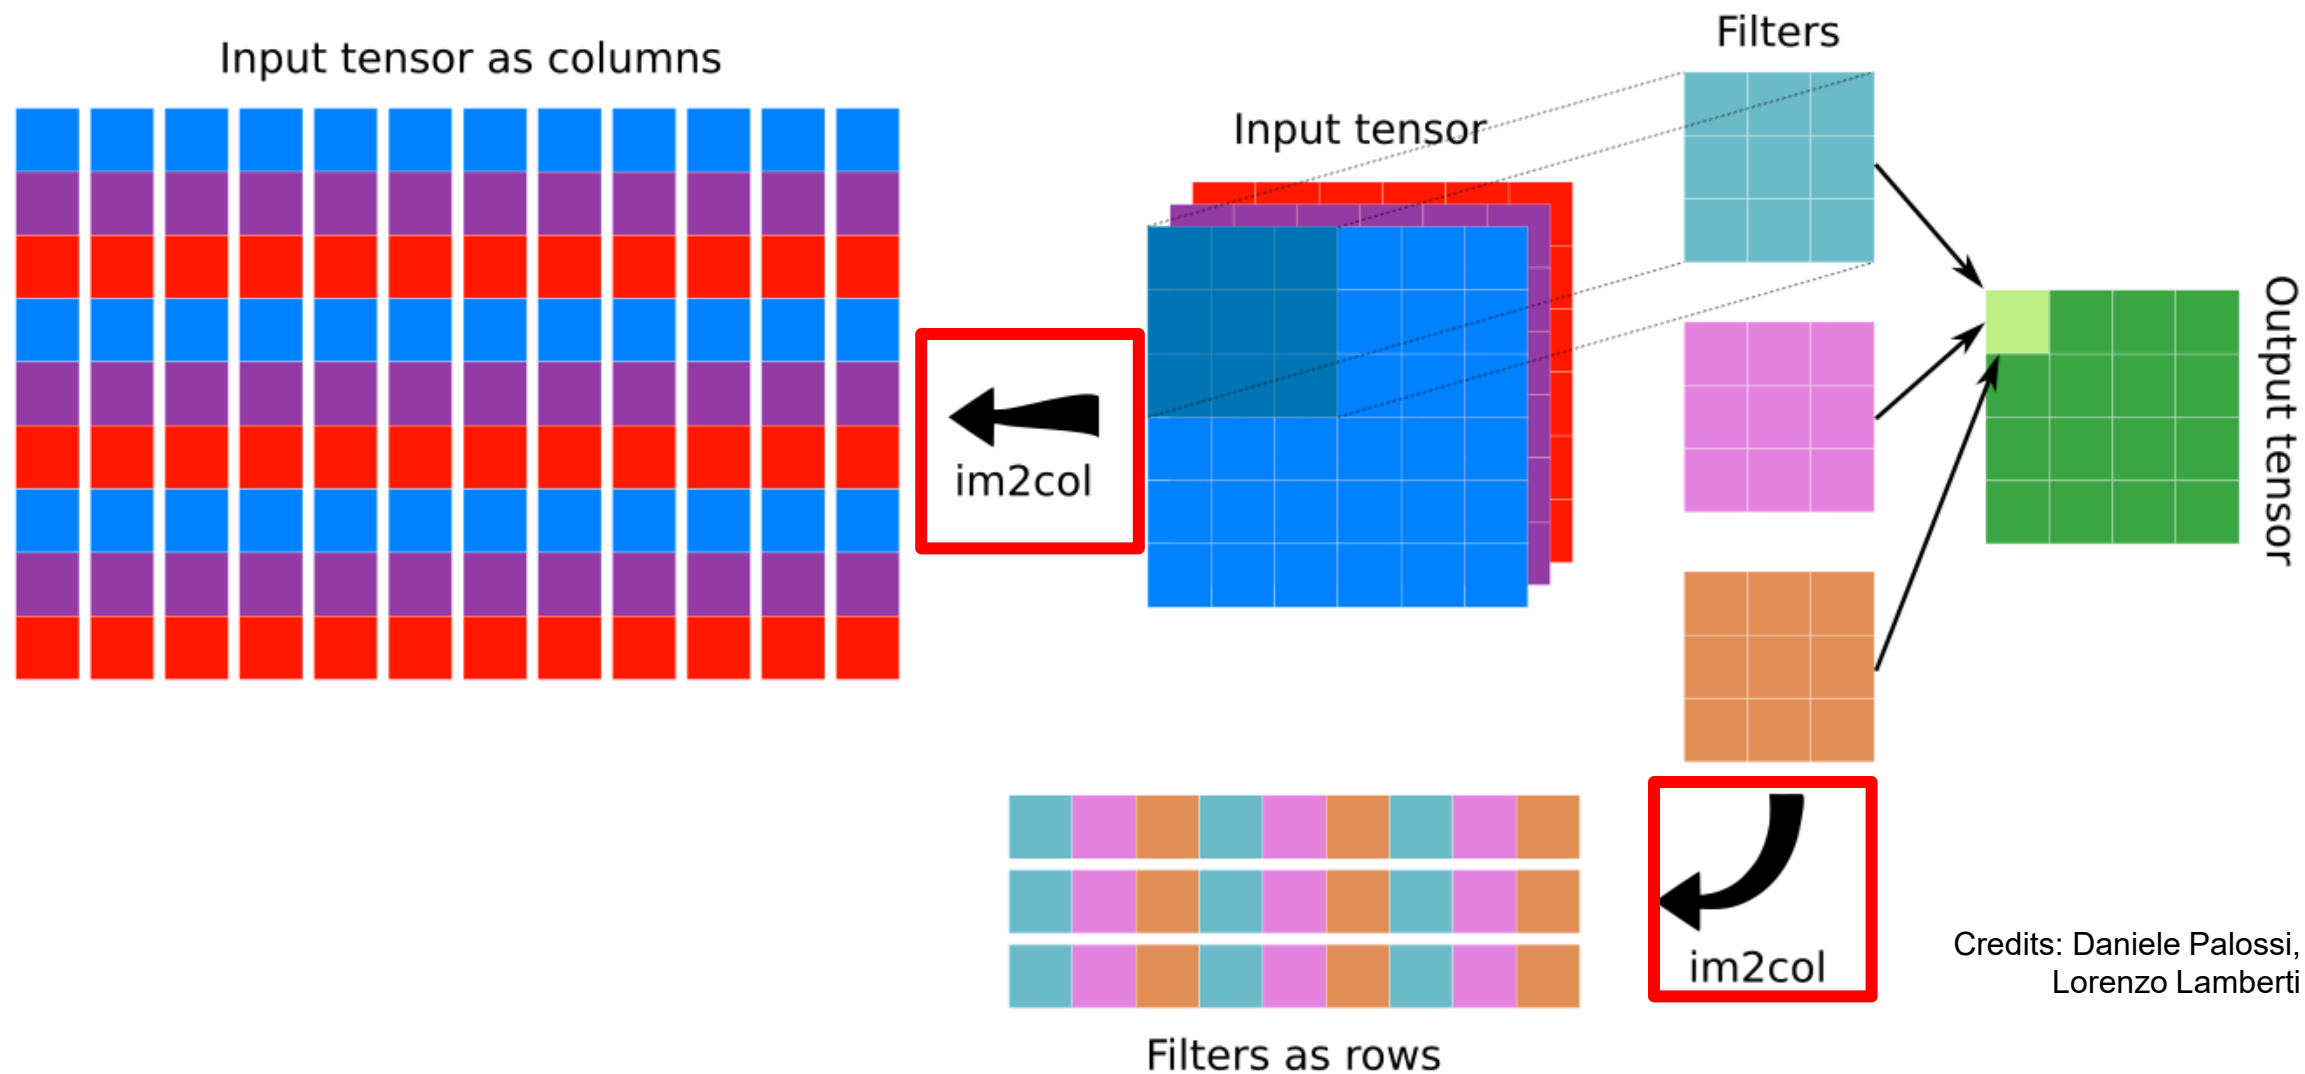
See example: http://bit.ly/hsdes21_conv2d

Convolution Operation: naive



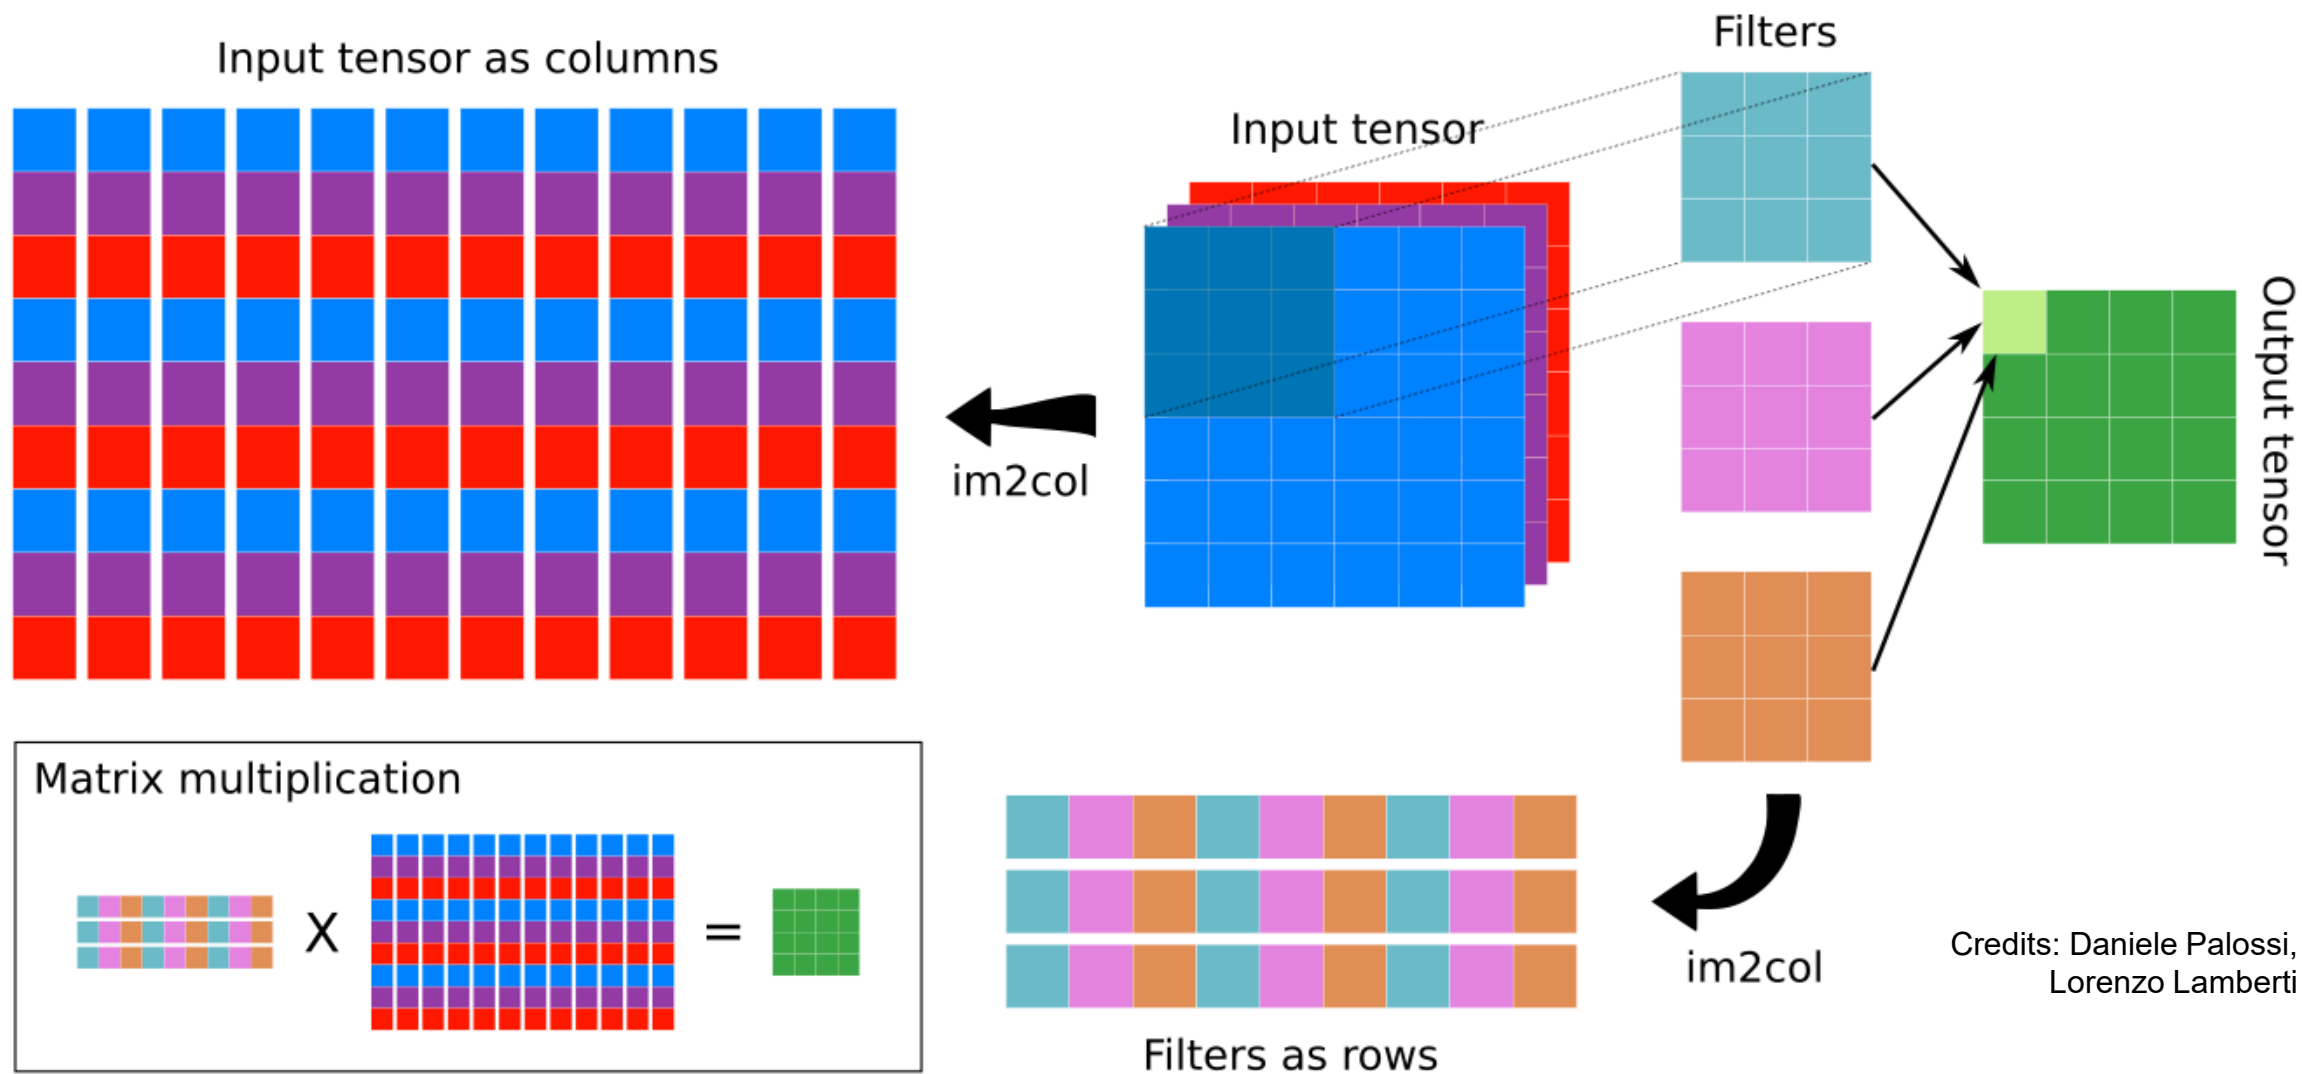
Credits: Daniele Palossi,
Lorenzo Lamberti

Convolution Operation: im2col and MatMul



Credits: Daniele Palossi,
Lorenzo Lamberti

Convolution Operation: im2col and MatMul



Credits: Daniele Palossi,
Lorenzo Lamberti

DNN Conv2D Layer on PULP (PULP-NN)

```
$ cd conv2d/  
$ make clean all run CORES=<1 to 8>
```

```
for i in range(0, H_out):  
    for j in range(0, W_out):  
        for m in range(0, K_out):  
            for ui in range(0, F):  
                for uj in range(0, F):  
                    for n in range(0, K_in):  
                        im2col[ui*F*K_in + uj*K_in + n] =  
                            x[i+ui, j+uj, n]  
                    psum = 0  
                    for idx in range(0, F*F*K_in):  
                        psum += w[m, idx] * im2col[idx]  
                    y[i, j, m] = act(psum)
```

Parallel loop

im2col

GEMM

Implementation details:

- Standard **output stationary** loop nest
- **HWC** layout for activations, **CoHWCi** for weights
- 2 *im2col* buffer per core!!
- Unrolled MatMul

- Individuate parallelism, im2col and GEMM in the code
- Measure #clk_cyc and #instr if running conv2d on 1, 2, 4 and 8 cores
 - What is the percentage of workload of the GEMM?

Benchmarking the Conv2D Kernel

\$ make dis

```
1c0091ec <pulp_nn_matmul_u8_i8>:
.....
1c0092a8: 01c4c0fb    lp.setup    x1,s1,1c0092e0 <pulp_nn_matmul_u8_i8+0xf4>
1c0092ac: 0049a60b    p.lw       a2,4(s3!)
1c0092b0: 0049268b    p.lw       a3,4(s2!)
1c0092b4: 004ba88b    p.lw       a7,4(s7!)
1c0092b8: 004b280b    p.lw       a6,4(s6!)
1c0092bc: 004aa50b    p.lw       a0,4(s5!)
1c0092c0: 004a258b    p.lw       a1,4(s4!) # 4004 <pos_soc_event_callback+0x3bbc>
1c0092c4: a9161f57    pv.sdotusp.b t5,a2,a7
1c0092c8: a9061ed7    pv.sdotusp.b t4,a2,a6
1c0092cc: a8a61e57    pv.sdotusp.b t3,a2,a0
1c0092d0: a8b61357    pv.sdotusp.b t1,a2,a1
1c0092d4: a9169457    pv.sdotusp.b s0,a3,a7
1c0092d8: a90693d7    pv.sdotusp.b t2,a3,a6
1c0092dc: a8a692d7    pv.sdotusp.b t0,a3,a0
1c0092e0: a8b69fd7    pv.sdotusp.b t6,a3,a1
.....
```

From the assembly we expect 6 LD + 8 VMAC to compute 8x4 8 bit MAC

- 14 instr / 32 MAC = 0.437 instr / MAC

	clk	clk/MAC	Speed Up	instr	inst/MAC
1 core	2335846	0.49503	1	2297016	0.486801
2 cores	1169969	0.247949	1.996502	1148625	0.243425
4 cores	588218	0.12466	3.971055	575739	0.122015
8 cores	298860	0.063337	7.815854	289294	0.061309

From the measurements we see:

- 1 core almost as expected (workload dominated by the matmul inner loop)
- #instr ~ #cyc -> no stall!!
- Almost linear speedup!!



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

DEI – Università di Bologna

www.unibo.it